

UNIVERSITÉ LIBRE DE BRUXELLES



ÉCOLE  
**POLYTECHNIQUE**  
DE BRUXELLES

INFO-H417 : DATABASE SYSTEMS ARCHITECTURE

---

Project report

---

*Professor:*

Stijn VANSUMMEREN

*Students :*

Miró-Manuel MATAGNE

Duc Minh NGUYEN

Linh TRAN NGOC

Academic Year 2020 - 2021

# Contents

|          |   |           |
|----------|---|-----------|
| <b>1</b> | <b>Introduction</b>   | <b>2</b>  |
| <b>2</b> | <b>Hardware and general setup</b>                             | <b>2</b>  |
| <b>3</b> | <b>Code architecture</b>                                      | <b>2</b>  |
| 3.1      | Input streams . . . . .                                       | 2         |
| 3.2      | Output streams . . . . .                                      | 2         |
| 3.3      | Experiments . . . . .   | 3         |
| 3.4      | Multi-way merge sort . . . . .                                | 3         |
| 3.5      | Testing . . . . .   | 3         |
| <b>4</b> | <b>Mapping</b>  | <b>3</b>  |
| <b>5</b> | <b>Experiment 1 : Sequential reading</b>                      | <b>4</b>  |
| 5.1      | First implementation . . . . .                                | 5         |
| 5.2      | Second implementation . . . . .                               | 6         |
| 5.3      | Third implementation . . . . .                                | 8         |
| 5.4      | Fourth implementation . . . . .                               | 11        |
| 5.5      | Comparison of the different implementations . . . . .         | 13        |
| <b>6</b> | <b>Experiment 2 : Random reading</b>                          | <b>16</b> |
| 6.1      | First implementation . . . . .                                | 17        |
| 6.2      | Second implementation . . . . .                               | 18        |
| 6.3      | Third implementation . . . . .                                | 20        |
| 6.4      | Fourth implementation . . . . .                               | 23        |
| 6.5      | Comparison of the different implementations . . . . .         | 26        |
| <b>7</b> | <b>Experiment 3 : Combined reading and writing</b>            | <b>27</b> |
| 7.1      | Variation of B . . . . .                                      | 29        |
| 7.2      | Comparison of the pairs of InputStream/OutputStream . . . . . | 30        |
| <b>8</b> | <b>Multi-way merge</b>  | <b>31</b> |
| <b>9</b> | <b>Conclusion</b>   | <b>35</b> |
| <b>A</b> | <b>Appendix 1 : Experiment 1</b>                              | <b>36</b> |
| A.1      | First implementation . . . . .                                | 36        |
| A.2      | Second implementation . . . . .                               | 37        |
| A.3      | Third implementation . . . . .                                | 38        |
| A.4      | Fourth implementation . . . . .                               | 40        |
| A.5      | Comparisons . . . . .   | 41        |
| <b>B</b> | <b>Appendix 2 : Experiment 2</b>                              | <b>42</b> |
| B.1      | First implementation . . . . .                                | 42        |
| B.2      | Second implementation . . . . .                               | 43        |
| B.3      | Third implementation . . . . .                                | 44        |
| B.4      | Fourth implementation . . . . .                               | 47        |
| <b>C</b> | <b>Appendix 3 : Experiment 3</b>                              | <b>49</b> |
| <b>D</b> | <b>Appendix 4 : Multi-way Merge sort</b>                      | <b>51</b> |

# 1 Introduction

The goal of this project, as part of the course of Database systems architecture, is to compare the performances of different methods for reading from and writing to secondary memory. The main source of performance loss is due to I/O operations, and we will therefore try to determine which implementations can reduce these I/O costs to the maximum, and which are less efficient. In addition, it is required to implement an external-memory merge-sort algorithm to examine its performances under different parameters.

In order to reach this goal, this report will describe first of all the machine environment on which the experiments were ran on. Then, after an explanation of the code architecture, the concept of mapping will be explained in detail. Finally, the different experiments as well as the Multi-way merge will be discussed.

## 2 Hardware and general setup

The implementation of this project was done in C++, using Visual Studio C++ on Windows. The virtual environment used was CLion for Windows.

All the measures presented in this report have been done on the same computer in order to be able to easily compare the values. This machine runs on Windows 10, uses a HDD disk, and has a 8GB memory. We are conscious that the results might be slower than on a machine using an SSD as access latency might be 1 or 2 orders of magnitude faster in that case.

## 3 Code architecture

The code architecture is generally speaking quite straightforward. The folders `InputStreams` and `OutputStreams` respectively group all the different input stream and output stream implementations. The folder `Experiments` groups all the 3 experiments that have been made throughout the project.

### 3.1 Input streams

All 4 input stream implementations inherit from one common parent class called `InputStream` in order to avoid duplicating code.

`InputStream1` uses the `read` function in order to read characters one by one until an end of line is detected. `InputStream2` makes use of `fgets` in order to fetch an entire line at the time from the considered file. `InputStream3` contains an internal buffer of size `B` (called `buffer`) and fills it up everytime it is required to read something that is not contained in it. `InputStream4` uses file mapping in order to read lines from a file. As this is an important part of our project, a theoretical point about file mappings will be made at section 4.

### 3.2 Output streams

All 4 output stream implementations also inherit from a common class called `OutputStream`.

These 4 implementations work very similarly to the input streams, in a symmetric way. `write` is used instead of `read` and `fputs` instead of `fgets`. Once again, `OutputStream4` makes use of file mappings.

### 3.3 Experiments

A separate class is made for each one of the experiments. Notice however that only the header files are present for these classes since all experiments implement a template function which allows to avoid repeating code for each different input or output classes we want to apply the experiment to. The template is able to execute the experiment for any input stream given for `Experiment1` and `Experiment2`, and for any combination of input and output streams for `Experiment3`.

### 3.4 Multi-way merge sort

The `MultiwayMerge` class is used to execute the external multi-way merge sort. It does not require a template since it uses only one input and output stream implementation.

### 3.5 Testing

The experiments have to be tested, and in particular their times of execution have to be measured in order to see the influence of each parameter on the efficiency of execution. The `Chrono` class is used in order to handle time measurements, it acts like a chronometer. All the tests on the several experiments as well as the Multi-way merge sort have been implemented in the `Measurement` class. Although this function might be quite long (there are many parameters to vary for each experiment), it is quite simple and is only used to measure execution times of experiments with varying parameters.

It is important to point out that all measurements present in this report are average values : the tests are executed multiple times in order to have more reliable values. The number of repetitions of a test depends on the expected execution time (when a test is extremely long, fewer repetitions will be made in order to avoid running tests for an excessive time).

## 4 Mapping

Since file mappings are present in our implementations and that it is not a very straightforward process, a theoretical explanation of this functionality is given.

Mapping a file can be compared to the action of creating a link between the disk (where the file is stored) and the physical memory (where the process' address space is). Indeed, the memory-mapped files allow us to access files on the disk as they were present in physical memory through pointers (like accessing dynamic memory). Reading a memory-mapped file (MMF) can be done by dereferencing a pointer in the designated range of addresses and similarly, writing data can be done by assigning a value to a dereferenced pointer. However, reading/writing to the disk is handled at a lower level to improve performance, so the file input/output (I/O) is not done exactly at the time the reading/writing of the memory-mapped file is performed.

MMF I/O operations have the advantage of being fast because data transfers are performed in a fix number of pages of data. Depending on the OS, the size of a page can vary. In our case, it is 4096 bytes and the MMF uses 16 pages (65536 bytes) for each I/O (it corresponds to the allocation granularity of the OS). Those pages of memory are managed by the virtual-memory manager (VMM) and because all the disk I/O operations have been standardised when it is performed by this VMM, it is a lot faster (routine memory accesses).

The main advantage of this method is that we can map a big file in memory even if this file is a lot bigger than the available RAM. Indeed, it is the principle of virtual memory that we use. We can access the whole file through a pointer as it was in RAM, but the I/O is performed only when accessing a part of the file not already in RAM. Each I/O contains 65 536 bytes (in our case) which obviously contains itself the part of the file needed. It is only if we want to read another part of

the file that is not in those 65 536 bytes that another I/O will be performed. Note that if multiple I/Os are performed, older chunks of data of the file will be removed from the RAM at some point (otherwise the whole concept of mapping would be irrelevant if we kept everything in RAM). So by using MMF I/O all the I/O interactions now occur in the user space (`read()` and `write()` include a pointer to a buffer in the process' space where the data is stored, so the kernel has to copy it, whereas MMF is directly in the file process' address space, so process has a direct access to the file).

An important remark has to be made nevertheless, using MMFs for simply reading a file into RAM compared to other methods of reading will not differ a lot in terms of performance. However, for different cases it can be a lot faster, if we want to modify a record of a large database by overwriting it, each time a new part of the record is needed, another file read is required (a lot of I/Os are needed). In the MMF approach, when the record is first accessed, all the page(s) of memory containing the record are read into memory. So, there is no system call overhead when accessing memory mapped files after the initial call. It may be possible that we do not see improvement of the performance using file mapping, here is why : methods that can use the disk cache will not be significantly slower unless a random reading is performed. Indeed, memory maps allow to keep using pages without additional overhead until you are done while `read()` does not.

Depending on systems, memory-mapped file functions differ. Here is how we implemented it on Windows with C++. To perform the memory mapping file, we call `CreateFileMapping`. This will only create a memory-mapped file object (very few resources needed even if it is a large file) into the physical memory. Then we need to map a view of a memory-mapped file, to do so, we need a valid handle to the MMF object (it is returned by the `CreateFileMapping` function). The `MapViewOfFile` function allows us to map a view of the file in order to perform our writing/reading. It is only when a view of the file is mapped that resources are loaded in the process's address space. We implemented this in such a way that we map and unmap parts of the file. Note, however, that only a multiple of the allocation granularity can be used for the offset (which is why we choose to round the size of the buffer B to a multiple of this value). In fact, this method of mapping part of the file is not really relevant (beside the pedagogical aspect), because one of the advantages of the view of the file is that we can read a file that is bigger than the process's address space, so every I/O is performed only when needed (the mapping does not imply I/O, only accessing a mapped region of the file will induce I/O !). Thus mapping the whole file will not bring the whole file in RAM, accessing part of the file mapped, however, will provoke an I/O (if the corresponding resources have not already been brought into RAM).

For the writing part of the project, we use a copy function in order to put data into the file (data and the file are both in the process's address space). It is like copying data from a buffer to another. Note that when the view is unmapped (or when the file-mapping object is deleted), changes are automatically written into the disk.

## 5 Experiment 1 : Sequential reading

The goal of this first experiment is to determine which of the 4 input stream implementations is more efficient for sequential reading, as well as finding the pros and cons of each implementation.

The code relative to this experiment is located in the `Experiment1` class, which implements a template function called `length`, that can be used with any input stream class. This function calls the `readln` function of the appropriate input stream class, and computes the length of each line, which is added to a counter `sum`. Once the end of the file is reached, the total length of the file is returned.

## 5.1 First implementation

In the first implementation of the input stream (class `InputStream1`), we used the `read` system call function, which can be seen as a function provided by the kernel. Without going into too much detail, a function from the system call is generally slower and need to switch between the user space and the kernel. The advantage is that we have more control and the reading is performed directly when asked (which is not the case for the `fgets` function for example). The `readln` function reads characters one by one until reaching the end of the line. Therefore, for each line, the number of necessary I/O operations to return the line is equal to the number of characters in the line. By extension, when reading the file completely and computing its length in this experiment, the total number of I/O operations is inevitably equal to the number of characters contained in the file, which is  $N$ . Therefore, we can write the cost function of `length` when the `InputStream1` class is used, which we will call `length1` as follows :

$$C(\text{length1}) = N \quad (1)$$

The number of I/O therefore evolves linearly in function of the size of the file being read. This relation can be verified by measuring the execution time of the function on several files of different sizes, and plotting the execution time in function of the file size. These measurements have been computed on every file of the dataset, taking the average execution time on 10 total executions per file. These results are available in the appendix A. The most meaningful results in order to verify the cost function are shown in the table 1 below.

| File name                   | File length (N) | Execution time in ms (t) | N/t    |
|-----------------------------|-----------------|--------------------------|--------|
| <code>aka_name</code>       | 73004383        | 288342,21                | 253,19 |
| <code>char_name</code>      | 215711567       | 868975                   | 248,24 |
| <code>comp_cast_type</code> | 45              | 1,93                     | 23,35  |
| <code>complete_cast</code>  | 2414495         | 9242,43                  | 261,24 |
| <code>movie_link</code>     | 656584          | 2611,33                  | 251,44 |

Table 1: Main results of Experiment 1 using `InputStream1`

As per the cost function defined at equation 1, the ratio  $\frac{N}{t}$  should be a constant. For most of the files tested, this value fluctuated between 245 and 265, but some values are particularly low. For instance, as shown in the table, the file `comp_cast_type` has a ratio of 23,35 which is surprisingly low compared to the rest of the results. This is justified by the very small length of that file. Indeed, in the case of very short files, the constant costs induced by the file openings and closings are not negligible for example. Therefore, the execution time will be higher when put in proportion to the file length. For larger files, the execution time is so big that the constant costs can be neglected, or at least their impact is way less noticeable. Notice that these surprisingly low ratio values occur for files that have a length up to around 2000 at least (the `info_type` file presents a ratio of 48, for a length of 1815, see appendix A).

By analyzing these results, or the complete table at appendix A.1 and neglecting the very short files, an average value for the ratio  $\frac{N}{t}$  is computed, and is equal to 246,74 ms. The standard deviation is then calculated on these same values, and equals to 17,94. In order to compare these values with the other implementations where the ratios will have much different values, we compute the coefficient of variation, defined as the quotient of the standard deviation by the mean. The results are summarized here below :

$$\begin{aligned}\mu &= 246,74 \\ \sigma &= 17,94 \\ c_v &= \frac{\sigma}{\mu} = 7,27\%\end{aligned}$$

The coefficient of variation is relatively small when the very short files are neglected, we can therefore consider the cost function established at equation 1 is correct. In order to confirm this result, a graph plotting the execution time of the `length` function on all files of the dataset is computed :

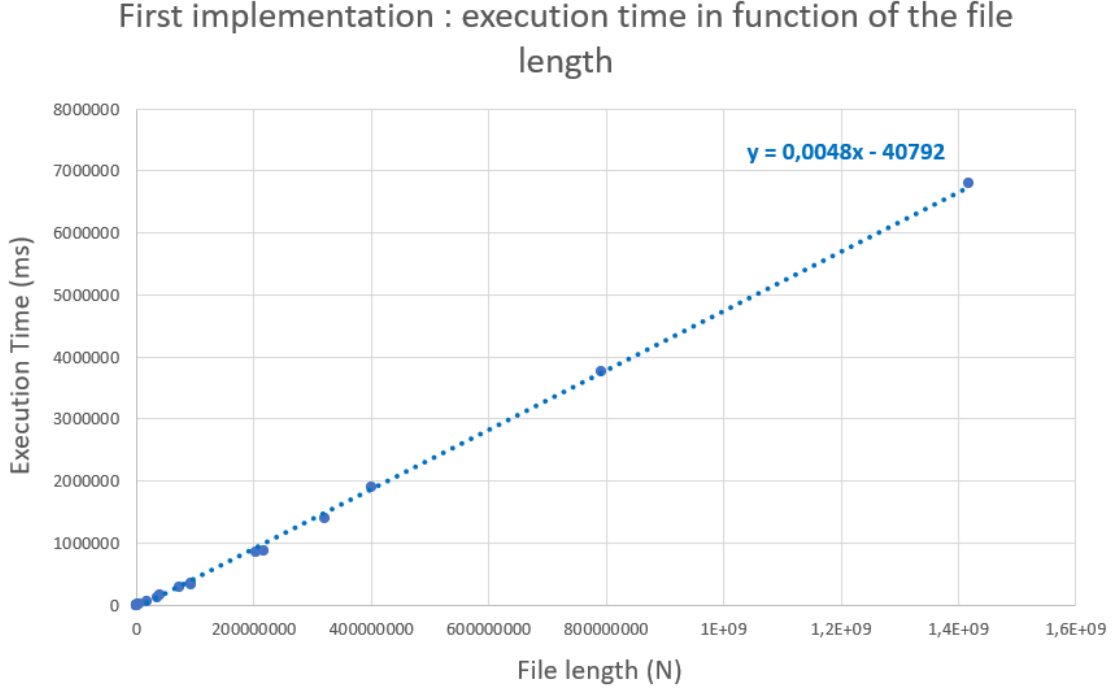


Figure 1: Execution time of the `length` function in Experiment 1 using `InputStream1` in function of the file length

The data seems to be arranged in a linear fashion. Indeed, the trendline of this set of data is computed and we obtain an affine function equation as displayed on the graph. Notice that the intercept is at 40792ms, which is very small compared to the execution times we measured.

Furthermore, we can estimate the cost of a single I/O operation involving the fetching of a single character. This is given by  $\frac{t}{N}$ , which gives, using the average computed previously,  $4,05 \cdot 10^{-6}s$ . Notice that this corresponds to the value of the slope of the trendline on the previous graph (where the values are in milliseconds). We were expecting a value quite a bit higher than this one (around  $10^{-3}s$  for a HDD for the access latency), we believe the fact that we are reading the file sequentially might considerably reduce the cost of the I/O since the disk head will always be right next to the next character to be read (if the data is stored sequentially on disk).

## 5.2 Second implementation

In the second implementation of the input stream, the `readln` function is based on the `fgets` function defined in the C library `stdio`. It is indeed a library call (so no mode switching like in a system call), and uses buffered I/O's which allows to use the cache to have better performance. This function uses an internal buffer in order to retrieve a line from a file. This function might be a little bit misleading, as it takes as argument an integer `n` representing the maximal number of

arguments to be read. One could therefore believe that this integer is the size of the buffer used to read a line from the file. However, by taking a look at the source code of `fgets`<sup>1</sup>, it is clear that the function used to refill the buffer (`__srefill`) does not depend on `n`. The calculations should therefore not be too different when modifying `n`, in terms of I/O operations. It is worth noticing that `fgets` is an optimized function that takes advantage of several hardware optimizations and uses the cache to improve performance. We will call `S` the size of the internal buffer, which is in our case equal to 4096 (as it is defined by the `BUFSIZ` constant of the `stdio` library). This buffer is refilled only when it is empty, therefore the cost function appears to be :

$$C(\text{length2}) = \left\lceil \frac{N}{S} \right\rceil \quad (2)$$

Notice the ceiling function around the fraction. This is explained by the fact that an I/O operation will be necessary to store the last elements of the file in a buffer, even though there might be less than `S` characters left in the file.

In order to verify this cost function, the execution time of this implementation of the Experiment 1 is computed on all the files of the dataset, and the results can be found in appendix A.4. The main findings, on files of different sizes are presented below :

| File name                   | File length (N) | Execution time in ms (t) | N/t      |
|-----------------------------|-----------------|--------------------------|----------|
| <code>aka_name</code>       | 73004383        | 3667,61                  | 19909,17 |
| <code>char_name</code>      | 215711567       | 11619,2                  | 18565,10 |
| <code>comp_cast_type</code> | 45              | 0,151                    | 298,01   |
| <code>complete_cast</code>  | 2414495         | 282,008                  | 8561,80  |
| <code>movie_link</code>     | 656584          | 69,1148                  | 9499,90  |

Table 2: Main results of Experiment 1 using `InputStream2`

First of all, it is clearly noticeable that this implementation is much faster than the first one. The  $\frac{N}{t}$  ratio seems to have values that are much more diversified than previously. This is due to the ceiling function, which breaks the linearity of the relation, and as a consequence leads to the fact that  $\frac{N}{t}$  is not a constant. Furthermore, this method is optimized in order to work rapidly with the operating system, the parameter `S` is chosen widely in order to optimize the efficiency of the I/O operations. As previously, results are quite different on the smaller files like `comp_cast_type.csv` due to the non negligible constant costs. By computing the average value of the  $\frac{N}{t}$  ratio, as well as the disparity of the values through the standard deviation and the coefficient of variation, we obtain the following results :

$$\begin{aligned} \mu &= 15109,41 \\ \sigma &= 4357,23 \\ c_v &= \frac{\sigma}{\mu} = 28,83\% \end{aligned}$$

The coefficient of variation is, as expected, higher than with the first implementation, due to the ceiling function present in the cost, and due to the optimizations of the `fgets` function which also might break the linearity. In order to confirm this cost function is correct, we can plot the execution times in function of the file lengths for all the files :

<sup>1</sup><https://android.googlesource.com/platform/bionic/+/ics-mr0/libc/stdio/fgets.c>



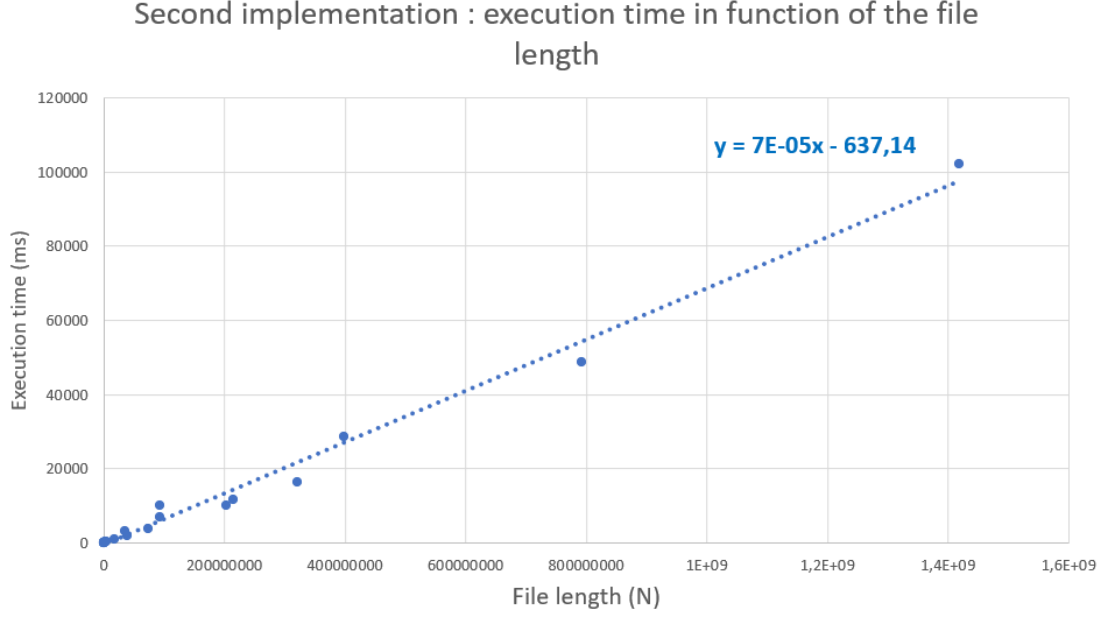


Figure 2: Execution time of the length function in Experiment 1 using `InputStream2` in function of the file length

We can observe that, even though the linearity is not as *strong* as in the previous implementation, the trendline still approximates the results relatively well, and that our cost function is probably correct for this implementation.

### 5.3 Third implementation

In the third implementation of the input stream (class `InputStream3`), we used the same reading function as in the first implementation but equipped with a buffer. Unlike the second implementation, no buffered I/O's are performed, only direct I/O's are being performed here. The `readln` function uses a buffer of size  $B$  in order to read a line from a file. Every time the buffer is empty, it is reloaded with the  $B$  next characters of the file to be read. Therefore, it is important to notice that there could be more than 1 I/O operation per line if  $B$  is smaller than the length of the considered line. On the other hand, there could also be less than 1 I/O operation per line, if the combined length of 2 successive lines is smaller than  $B$  for instance. Therefore, in order to read a file in its entirety to find the total length, the cost is directly related to the number of bufferings of  $B$  characters that will be necessary throughout. This allows us to define a cost function for the `length` function using the `InputStream3` class, which we will call `length3` :

$$C(\text{length3}) = \left\lceil \frac{N}{B} \right\rceil \quad (3)$$

This cost function is therefore inversely proportional to  $B$  when the same file is being treated. When  $B$  is maintained constant, the cost should evolve quite linearly in regard to the number of characters in the file, although the ceiling function should introduce *steps* in this function which slightly breaks the linearity. Table 26 shows an overview of the results obtained for  $B = 100$ . The complete results are in appendix A.3.

| File name      | File length (N) | Execution time (t) | ceil(N/B) | ceil(N/B)/t |
|----------------|-----------------|--------------------|-----------|-------------|
| aka_name       | 73004383        | 8785,92            | 7300439   | 830,92      |
| char_name      | 215711567       | 25043,9            | 21571157  | 861,33      |
| comp_cast_type | 45              | 14,50              | 5         | 0,34        |
| complete_cast  | 2414495         | 338,03             | 241450    | 714,28      |
| movie_link     | 656584          | 82,33              | 65659     | 797,49      |

Table 3: Results of Experiment 1 using InputStream3

This table outputs the ratio  $\frac{\lceil \frac{N}{B} \rceil}{t}$ , which should be close to a constant. We do however observe that for a constant B, not only is this value highly impacted by the small files (as before), but the fluctuations of the ratio also seem a bit more important than during the first implementation. By computing the average value of that ratio, as well as the standard deviation and the coefficient of variation without considering the very small files, we obtain :

$$\begin{aligned}\mu &= 811,42 \\ \sigma &= 109,81 \\ c_v &= \frac{\sigma}{\mu} = 13,53\%\end{aligned}$$

Although this coefficient of variation remains small enough in order to confirm the cost function established earlier for a fixed value of B, it is nearly twice higher than the one obtained during the first implementation. We suppose these variations are due to the fact that although the number of I/O operations is given by the formula 3, the duration of each ones of these I/O operations will not necessarily be the same. Indeed, the last I/O operation will typically be shorter : it has the same access latency as the previous ones, but less data has to be transferred on average, so the time taken for the throughput will typically be smaller. On top of that, the ceiling function breaks the linearity as described previously. Nevertheless, we can build a graph in order to prove the linear character of the cost function 3 with a constant B :




Figure 3: Execution time of the length function in Experiment 1 using InputStream3 with a constant B = 100, in function of the file length

We see the linear trendline approximates quite well the values plotted, and has an intercept of 988.98ms, which is very small compared to the measured lengths. Therefore, these results obtained

with  $B = 100$  are a first justification of the cost formula 3.

Next, it is necessary to show the impact of  $B$  on the execution time. In order to do that, a same file `aka_name` is read multiple times with several different  $B$  values. The main findings of this experiment are presented below :

| B   | Execution time is ms (t) | ceil(N/B) | ceil(N/B)/t |
|-----|--------------------------|-----------|-------------|
| 10  | 31596,1                  | 7300439   | 231,06      |
| 50  | 10836,9                  | 1460088   | 134,73      |
| 100 | 8151,63                  | 730044    | 89,56       |
| 150 | 7225,41                  | 486696    | 67,36       |
| 200 | 6692,75                  | 365022    | 54,54       |

Table 4: Main results of Experiment 1 using `InputStream3` on `aka_name`

We do indeed observe that when  $B$  gets bigger, the execution time tends to decrease. Furthermore, the hypothesis we made concerning the variability of the ratio  $\frac{\lceil \frac{N}{B} \rceil}{t}$  is way more visible in this case and seems to be correct : when  $B$  is bigger, the cost of each I/O operation is bigger (so the ratio decreases when  $B$  increases) because there is more data to be fetched. This is therefore a limitation of our cost function, which represents the number of I/O operations, but is not completely and directly linked to the execution time. By plotting the results of this experiment on  $B$  values going from 10 to 1000 with a step of 10 on the `aka_name` file, the following graph is obtained :




Figure 4: Execution time of the length function in Experiment 1 using `InputStream3` on the `aka_name` file in function of  $B$

The full results are displayed at Appendix A.3.

We do indeed see decreasing times when  $B$  increases, which confirms the established cost function. This method however presents some limitations, in particular from around  $B = 500$ , where the execution time seems to become constant no matter  $B$ . Concerning the execution times observed, they are greater than the ones observed on the same file with the second implementation using `fgets`.

## 5.4 Fourth implementation

For the fourth implementation, we used file mapping to read. For a detailed explanation of this method, please refer to section 4. As the file is read sequentially, the cost function for the fourth implementation is the following, where  $G$  is the allocation granularity (explained in section 4) :

$$C(\text{length4}) = \left\lceil \frac{N}{G} \right\rceil \quad (4)$$

$G$  is equal to 65536 ( $2^{16}$ ) on the machine the tests have been ran on.

Indeed, as a sequential reading is performed,  $B$  has no influence on the cost because whatever size of the file we map, we will never want to access information in pages that were fetched previously during the execution. The number of I/O operations performed is therefore directly related to the allocation granularity.

First, execution times of the `length` function using memory mapping on several files of different sizes are computed, in order to prove the pseudo-linear characteristic of the cost function established at equation 4 in function of  $N$ . These measurements are done with a constant value of  $B = 65536$ . The full results are displayed at Appendix A.4, and some of the main findings are reported in the table below :

| File name                   | File length (N) | Execution time (t) | ceil(N/G) | ceil(N/G)/t |
|-----------------------------|-----------------|--------------------|-----------|-------------|
| <code>aka_name</code>       | 73004383        | 5361,78            | 1140      | 0,21        |
| <code>char_name</code>      | 215711567       | 16252,8            | 3292      | 0,20        |
| <code>comp_cast_type</code> | 45              | 0,25               | 5         | 3,99        |
| <code>complete_cast</code>  | 2414495         | 317,221            | 37        | 0,12        |
| <code>movie_link</code>     | 656584          | 85,5874            | 11        | 0,13        |

Table 5: Results of Experiment 1 using `InputStream4`, with  $B = 65536$

From these measurements, we can expect  $\frac{\lceil \frac{N}{G} \rceil}{t}$  to be a constant. As previously explained, the shorter files present dissimilarities compared to the other ones because of the important constant costs. We can therefore compute the average value of this ratio, as well as the standard deviation and the coefficient of variation in order to evaluate if it is relatively close to being constant. These measurements do not include the files that are considered too small which give non-interpretable data, the results are shown below :

$$\begin{aligned} \mu &= 0,192 \\ \sigma &= 0,042 \\ c_v &= \frac{\sigma}{\mu} = 21,8\% \end{aligned}$$

The coefficient of variation is higher than the one we got during the first implementation, but is still low enough in order to consider this data is pseudo-constant. We can also point out that this solution seems quite comparable with the third implementation in terms of execution time (when  $B$  is very large), but also less efficient than the second implementation using `fgets`. By plotting the results of this experiment, we obtain the following graph :




Figure 5: Execution time of the length function in Experiment 1 using InputStream4 for  $B=65536$

We can therefore confirm the previously obtained results : the relation seems linear, but the fluctuations around the trendline seem higher than they were for the 2 first implementations.

In order to justify completely the cost function established at equation 4, we must prove that  $B$  does not influence the execution time. In order to do that, the function was applied to a single file `aka_title` with  $B$  values ranging from 65536 to 3276800 (50 times  $G$ ). The results are briefly described in the table below, as well as in the following graph :

| B       | Execution time is ms (t) |
|---------|--------------------------|
| 131072  | 2369,42                  |
| 524288  | 2318,54                  |
| 1048576 | 2355,03                  |
| 1703936 | 2366,81                  |
| 2752512 | 2346,66                  |

Table 6: Main results of Experiment 1 using InputStream4 on `aka_title`




Figure 6: Execution time of the length function in Experiment 1 using InputStream4 on aka\_title in function of B

We can indeed see that the execution time seems to be constant whatever B. This can be proven by calculating, as previously, the mean, standard deviation and coefficient of variation of this list of execution times :

$$\begin{aligned}\mu &= 2362,71 \\ \sigma &= 33,22 \\ c_v &= \frac{\sigma}{\mu} = 1,4\%\end{aligned}$$

This extremely small coefficient of variation indeed indicates that the execution time is independent of B as expected.

## 5.5 Comparison of the different implementations

By simply analyzing the data we have collected up to now, it seems obvious that the sequential reading of a file using the first implementation of the `readln` function is definitely not the best. This result is expected, as the number of I/O operations is equal to the number of characters in the file, which is technically the maximum number of operations to do in order to read a file (without any unnecessary operations).

By considering the established cost functions, one could think the third implementation would work best, when a huge value of B is used. This is however not the case, because the times tend to stagnate quite rapidly when B increases as it was shown on figure 4.

By plotting all the execution times for the 4 implementations in function of B for the `aka_name` file, we obtain the following results :




Figure 7: Execution time of the length function in Experiment 1 for all different implementations on aka\_name

It is now clear that the first implementation is less efficient than the 3 others, so it is removed from the graph in order to allow clearer interpretation :

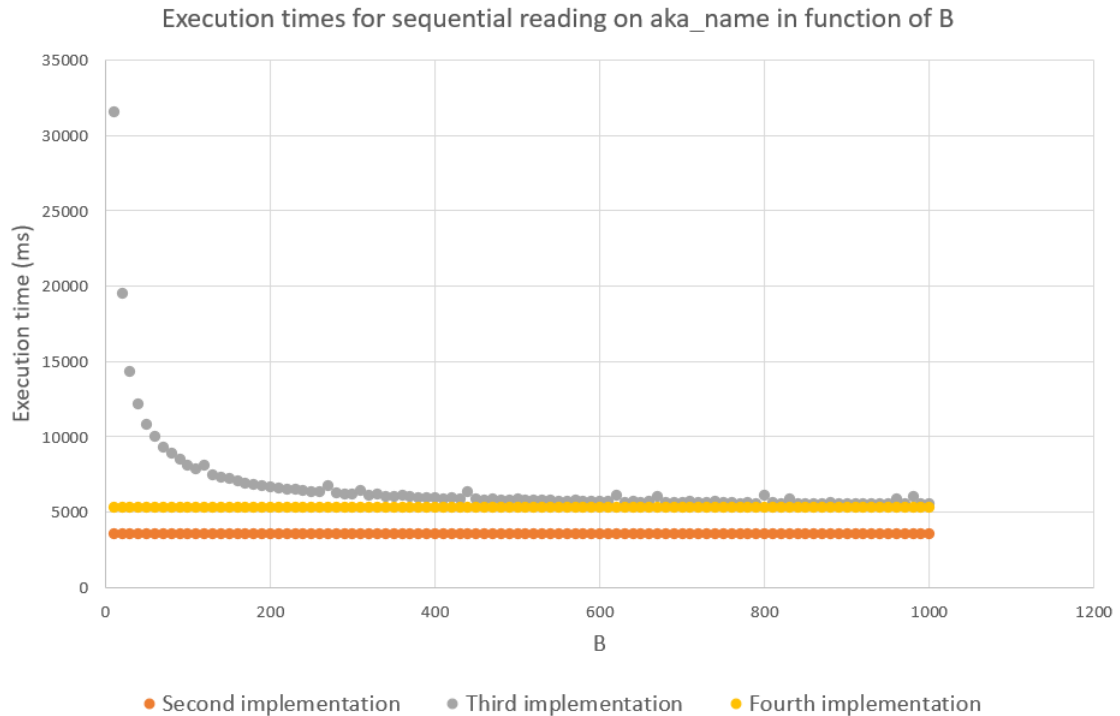


Figure 8: Execution time of the length function in Experiment 1 for the second, third and fourth implementations on aka\_name

The second implementation, using the built-in `fgets` method, is therefore the most efficient to execute sequential reading on `aka_name`, which has around  $7,3 \cdot 10^7$  characters. A justification for this behavior would be that `fgets` uses the computer cache very efficiently. Indeed, this function is

extremely optimized on order to adapt to the hardware, and the buffer sizes are chosen specially to allow the usage of cache, which is extremely rapid. The third implementation using the buffering is comparable to the memory mapping in terms of time when  $B$  is very large. The fourth implementation using the mapping is not particularly efficient in the case of sequential reading, because the advantage of mapping is mainly visible when we want to access data that has already been previously read.

We observe that time taken by the direct I/O's performed by the system call function is longer than the time taken by the buffered I/O's and the file mapping. Moreover, we can see that for this specific case (sequential reading) the buffered I/O's (managed by the OS) performed better than the file mapping. As explained above, we can deduce that the usage of cache is very well optimized for this kind of reading. These results must be validated on other files of various sizes. The same graph is computed for `complete_cast` (size  $2.10^6$ ) at figure 9 and for a smaller file `movie_link` (size  $6.10^5$ ) at figure 10.

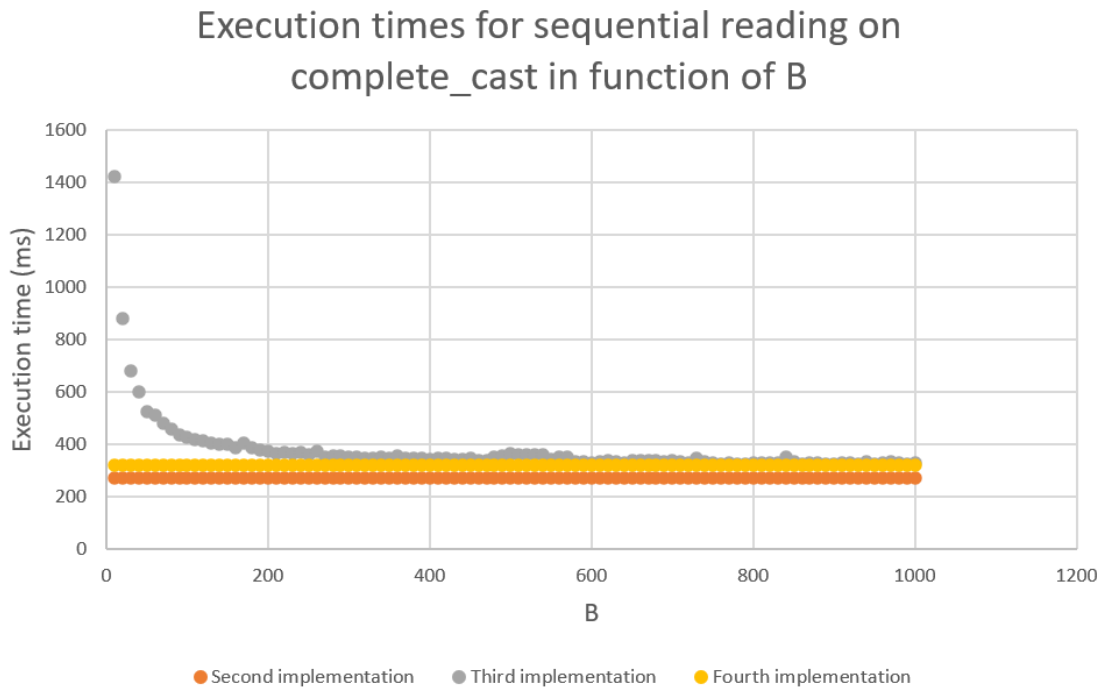


Figure 9: Execution time of the length function in Experiment 1 for the second, third and fourth implementations on `complete_cast`






Figure 10: Execution time of the length function in Experiment 1 for the second, third and fourth implementations on `movie_link`

The results obtained are very similar to the previous ones, the second implementation always giving slightly better results, and the third and fourth implementation giving extremely similar results when `B` is very large. It is important to remember the extremely short files are neglected in this analysis, because the constant costs arising from the file openings and closings for example are non negligible and therefore don't really give useful data to study the I/O costs.

## 6 Experiment 2 : Random reading

The objective of this second experiment is to determine which of the 4 input stream implementations is more efficient for random reading, as well as finding the pros and cons of each implementation.

The code relative to this experiment is located in the `Experiment2` class, which implements a template function called `randjump`, as the `Experiment1` class, that can be used with any input stream class. This function calls the `readln` function of the appropriate input stream class, but in contrary to the first experiment, it takes as input a parameter `j` which indicates the number of times the following process will be repeated :

- compute a random number position between 0 and the size of the input file;
- seek to the computed position, i.e. set the cursor to the computed position;
- add the length of the string returned by `readln` to a counter `sum`.

Once the process above is done `j` times, the total of the counter `sum` is returned.

Before the actual measurements on the 4 different implementations, it was expected that random reading would bring some differences compared to the first experiment through the way it reads the files. Indeed, in the first experiment, the implementations are compared by observing their execution times in function of the size of the files, while the random reading might be independent to the size of the files due to the fact that it will not read entirely a file from the start to the end but by jumping to different positions randomly generated from a seed which will be a constant in order to to have the same scenarios for the 4 `randjump` implementations. Therefore, the execution times of a random reading will obviously be dependent on the number of iterations `j` given in

parameter. It is expected that the second implementation of `InputStream` will be less efficient than previously since it uses an internal buffer, which will need to be refilled many times in the case of random readings. The same reasoning applies to the 3rd implementation, where the difference should be even more noticeable, because no cache optimization is used for this one. The fourth implementation should be way more efficient here, as the advantages of file mappings fully come in hand when randomly reading a file.

## 6.1 First implementation

As mentioned in the subsection 5.1, in the first implementation of the input stream (class `InputStream1`), the `readln` function reads characters one by one until reaching the end of the line and as a result, the number of necessary I/O operations to return a line is equal to the number of characters in each line. However, the `randjump` function makes  $j$  random jumps in the file, it means that `readln` will not be necessary called from the start of a line. As a result, the total number of I/O operations is inevitably equal to the product of 1 (length of the string returned by `readln`) and  $j$  (number of jumps in the file). Therefore, we can write the cost function of `randjump` when the `InputStream1` class is used, which we will call `randjump1` as follows :

$$C(\text{randjump1}) = l * j \quad (5)$$

The number of I/O therefore evolves in function of 2 parameters,  $l$  and  $j$ . As a result, two measurements are made, the first one is the execution time of the function on several files of different sizes by plotting the execution time in function of the file size with a fixed number of iterations  $j$  as a parameter. Then, the second measurement is the execution time of `randjump1` in function of the number of iterations with an arbitrary file.

The first measurement has been computed on multiple files of the dataset with a seed = 10 and a seed = 100 as shown in the tables 7 and 8 below. As the first experiment, the function `testFiles2` from the `Measurement` class is used by taking the average execution time on 10 total executions per file with a constant number of iterations  $j = 10$ .

| File name      | Execution time in ms (t) | Sum |
|----------------|--------------------------|-----|
| cast_info      | 15.0057                  | 273 |
| comp_cast_type | 1.99667                  | 82  |
| keyword        | 6.99008                  | 141 |
| movie_info_idx | 3.92387                  | 117 |

Table 7: The sums returned by the `randjump` function in Experiment 2 using `InputStream1` with a constant  $j = 10$  and a seed = 10

| File name      | Execution time in ms (t) | Sum |
|----------------|--------------------------|-----|
| cast_info      | 12.6152                  | 257 |
| comp_cast_type | 0.68175                  | 62  |
| keyword        | 6.21133                  | 132 |
| movie_info_idx | 9.60508                  | 136 |

Table 8: The sums returned by the `randjump` function in Experiment 2 using `InputStream1` with a constant  $j = 10$  and a seed = 100

The biggest file `cast_info`, in term of the total size, has the largest sum due to the longer size of the lines. However, it is neither proportional to the total size of a file nor the size of the lines because of the randomness, but it has an influence on the sum output. Furthermore, the randomness and the non-proportionality can be proven by noticing that with a seed equal to 10, the sum returned in the file `keyword` is bigger than the one with the file `movie_info_idx`. One would be tempted to say that it is due to the longer lines in `keyword` but with a constant seed equal to 100 it is the inverse, i.e. the sum returned with the file `movie_info_idx` is higher than the one returned with `keyword`. We can notice that the execution time evolves in regard to the sum.

In the second measurement, the `randjump1` function reads the file `cast_info` by varying the number of iterations `j`. As before, the average execution time on 10 total executions is taken. It is obvious that more the number of iterations `j` increases, and more the sum returned by `randjump1` will increase as well as the execution time. In order to confirm this statement, a graph plotting the execution time of the `randjump` function on the iterations applied on the file `cast_info` is computed :




Figure 11: Execution time of the `randjump` function in Experiment 2 using `InputStream1` on the `cast_info` file in function of the number of iterations

The graph seems to be arranged in a linear fashion despite some weak variations. Indeed, the trendline of this set of data is computed and an affine function equation is obtained as displayed on the graph.

The same linear behavior is noticed for larger numbers of iterations as shown in the table 30 of the appendix B.1.

## 6.2 Second implementation

Unlike the second implementation in the first experiment, using the second implementation in random reading is less efficient due to the fact that the size of the internal buffer (called `S`) of the `fgets` function does not have influence anymore. Indeed, the buffer will have to be refilled multiple

times for each call of the `randjump` function.

In order to observe the `randjump` function using the `InputStream2` class which we will call `randjump2`, 2 measurements will be computed, as the previous subsection. The first measurement computes the execution time of `randjump2` in function of the total size of different files of the dataset (taking the average execution time on 10 total executions per file). Then the second one computes the execution time of `randjump2` by varying the number of iterations  $j$ .

As shown on the two tables 9 and 10 below, the exact same observation than the first implementation can be made by running `randjump2` on the same files as the first experiment to put in evidence the differences. It confirms that the total size of a file does not influence the sum returned by the `randjump2` function. One difference we can notice is that the execution time is faster than the `randjump1`.

| File name                   | Execution time in ms (t) | Sum |
|-----------------------------|--------------------------|-----|
| <code>cast_info</code>      | 0.51673                  | 273 |
| <code>comp_cast_type</code> | 0.37535                  | 82  |
| <code>keyword</code>        | 0.33984                  | 141 |
| <code>movie_info_idx</code> | 0.37693                  | 117 |

Table 9: The sums returned by the `randjump` function in Experiment 2 using `InputStream2` with a constant  $j = 10$  and a seed = 10

| File name                   | Execution time in ms (t) | Sum |
|-----------------------------|--------------------------|-----|
| <code>cast_info</code>      | 9,97019                  | 257 |
| <code>comp_cast_type</code> | 0,57382                  | 62  |
| <code>keyword</code>        | 5,39875                  | 132 |
| <code>movie_info_idx</code> | 5,35759                  | 136 |

Table 10: The sums returned by the `randjump` function in Experiment 2 using `InputStream2` with a constant  $j = 10$  and a seed = 100

In the second measurement, the `randjump2` function reads the file `cast_info` by varying the number of iterations  $j$ . As the first implementation, more the number of iterations  $j$  increases, and more the sum returned by `randjump2` will increase as well as the execution time. In order to confirm this statement, a graph plotting the execution time of the `randjump` function on the iterations applied on the file `cast_info` is computed :




Figure 12: Execution time of the `randjump` function in Experiment 2 using `InputStream2` on the `cast_info` file in function of the number of iterations

The graph seems linear with a few variations. The results for larger numbers of iterations are displayed in the appendix B.2 and feature a similar linear fashion. Notice that `randjump2` is faster than `randjump1` as shown on the figure 12.

### 6.3 Third implementation

In this third implementation of the `InputStream3` class, the `readln` function uses a buffer of size `B` in order to read a line from a file as already said in the first experiment. Every time the buffer is empty, it is reloaded with the `B` next characters of the file to be read. However, since it is random reading, random jumps occur within the file being read, for each jump, the buffer is refilled. Therefore, as for the second implementation, its performance will be reduced in comparison to the first experiment. Note that the `randjump` function using the `InputStream3` class will be called `randjump3` from now on to simplify.

Three measurements are made, the first one is the execution time of the function on several files of different sizes by plotting the execution time in function of the size file with a number fixed of iterations `j` as a parameter and a constant `B` equal to 100. Then the second measurement is the execution time of `randjump3` in function of the number of iterations with an arbitrary file (constant `B = 100`). And finally, the execution time is plotted in function of different values of `B`.

As the previous implementations, the first measurement has been computed on multiple files of the dataset with a seed = 10 and a seed = 100 to prove the non-influence of the total size of a file, as shown in the tables 11 and 12 as follows :

| File name      | Execution time in ms (t) | Sum |
|----------------|--------------------------|-----|
| cast_info      | 0.49542                  | 273 |
| comp_cast_type | 1.67674                  | 82  |
| keyword        | 3.16489                  | 141 |
| movie_info_idx | 7.68705                  | 117 |

Table 11: The sums returned by the randjump function in Experiment 2 using InputStream3 with a constant  $j = 10$ ,  $B = 100$  and a seed = 10

| File name      | Execution time in ms (t) | Sum |
|----------------|--------------------------|-----|
| cast_info      | 0.32084                  | 257 |
| comp_cast_type | 0.3616                   | 62  |
| keyword        | 0.55428                  | 132 |
| movie_info_idx | 0.4215                   | 136 |

Table 12: The sums returned by the randjump function in Experiment 2 using InputStream3 with a constant  $j = 10$ ,  $B = 100$  and a seed = 100

Like the two previous implementations, we can see that the sum returned from the **keyword** file is greater than the one returned from the **movie\_info\_idx** file for a seed equal to 10 but for a seed equal to 100, the inverse occurs. It shows therefore the sums returned are indeed not proportional to the total size of a file.

In the second measurement, the **randjump3** function reads the file **cast\_info** by varying the number of iterations  $j$  with a fixed buffer size  $B$  equal to 100. As before, the average execution time on 10 total executions is taken. We will plot the execution time in function of the first 1 to 100 iterations :




Figure 13: Execution time of the randjump function in Experiment 2 using InputStream3 on the cast\_info file with a constant  $B = 100$ , in function of the number of iterations

The graph seems not linear and the execution times are very low and random, we suppose that these random variations are due to the low value of iterations of the **randjump3** which is executed too rapidly, and because of the constant costs such as the **seek**, **open** and **ftell** functions. In order to observe more accurately the **randjump3** function, we will plot on larger iterations as 1000

to 10000 as shown on the Figure 14 below :




Figure 14: Execution time of the randjump function in Experiment 2 using InputStream3 on the cast\_info file with a constant  $B = 100$ , in function of the number of iterations

Now we see the linear trendline approximates quite well the values plotted.

Next, it is necessary to show the impact of  $B$  on the execution time. In order to do so, a same file `cast_info` is read multiple times with a number of iterations  $j$  equal to 10000 with several different  $B$  values. The main findings of this experiment are presented below :




Figure 15: Execution time of the `randjump` function in Experiment 2 using `InputStream3` on the `cast_info` file in function of `B` with `j = 10000`

The full results are displayed in the tables 33 and 34 of the Appendix B.3.

In contrary to the first experiment, the execution time in function of `B` grows linearly, i.e. when `B` gets bigger, the execution time tends to increase. As a result, when `B` is bigger, the cost of each I/O operation is bigger because there is more data to be fetched, thus, the time taken for the throughput will be greater. However, since in the `randjump3` function, the buffer of size `B` is reset on each call of the `seek` function, the previous buffer is never used. Therefore, it is useless to increase `B` to very large values in this implementation, and it leads to large execution times.

#### 6.4 Fourth implementation

For this last implementation of random reading, 3 measurements will be computed as for the previous implementation. We will call the `randjump` function using the `InputStream4` class `randjump4`. In contrary to the sequential reading where `B` had no influence on the execution time, it will be important in this case because as the reading is random, we might need to access pages that were already fetched into memory previously. Therefore, the size `B` will have a direct influence on the execution time of the `randjump4`.

The first measurement is made once again to prove the non influence of the total size of a file on the sequential reading. This measurement is done on multiple files of the dataset with a constant number of iterations `j` and a fixed `B = 65536` (allocation granularity) and we vary the seed between 10 and 100. The results are the same as the others implementations as shown on the 2 tables 13 and 14, i.e. `randjump4` does not depend on the total size of a file.



| File name      | Execution time in ms (t) | Sum |
|----------------|--------------------------|-----|
| cast_info      | 4,06032                  | 273 |
| comp_cast_type | 0,79132                  | 82  |
| keyword        | 6,59532                  | 141 |
| movie_info_idx | 9,16791                  | 117 |

Table 13: The sums returned by the randjump function in Experiment 2 using InputStream4 with a constant  $j = 10$ ,  $B = 655336$  and a seed = 10

| File name      | Execution time in ms (t) | Sum |
|----------------|--------------------------|-----|
| cast_info      | 13,7111                  | 257 |
| comp_cast_type | 0,26802                  | 62  |
| keyword        | 7,94977                  | 132 |
| movie_info_idx | 13,1686                  | 136 |

Table 14: The sums returned by the randjump function in Experiment 2 using InputStream4 with a constant  $j = 10$ ,  $B = 655336$  and a seed = 100

Then the second measurement will be computed by plotting the **randjump4** on the **cast\_info** file with a constant  $B$  equal to the allocation granularity in function of the number of iterations as shown on the Figure 16 below :




Figure 16: Execution time of the randjump function in Experiment 2 using InputStream4 on the cast\_info file with a constant  $B = 65536$ , in function of the number of iterations

The graph features a linear behavior with some fluctuations as expected. Indeed, the execution time will be more random due to the small values of  $j$  and the constant costs of other function in the **randjump4** as for the third implementation. The execution on larger number of iterations presents a more linear behavior (see the table 35 in the Appendix B.4). Notice that the results are not so different in comparison to the third implementation in terms of execution time when  $j$  is very large, but less efficient than the second implementation.

Finally, in order to completely observe the **randjump** function using the **InputStream4** class,

we will plot the execution time of this function on the `cast_info` file by varying the size  $B$  from 655360 ( $10 \times 65536$ ) to 65536000 ( $1000 \times 65536$ ). The results are shown in the following graph :




Figure 17: Execution time of the `randjump` function in Experiment 2 using `InputStream4` on the `cast_info` file with a constant  $j = 10000$ , in function of  $B$  varying from 655360 to 65536000

The complete results are displayed in the table 36 in the Appendix B.4.

As  $B$  grows, we see that the execution time is relatively constant despite some variations. It is due to the fact that it is random. Indeed, with a different seed the graph will be a bit different dependently to the fact that whether the positions generated will be in the previously fetched pages or not. The execution time will therefore be different (when the position generated is not in the current mapped part, it will be unmapped, and we map the page in which the current position is).

However, when  $B$  is at bigger values, the execution time decreases drastically. We can explain it by the fact that when  $B$  is big, the part of the file being mapped will consequently grow. As a result, if the part being mapped gets bigger, the probability of a generated position being in the already mapped page is higher, i.e. less unmappings will occur and the execution time will be smaller. Indeed, unmapping parts of a file that have already been accessed will make the previous I/O unusable and so, due to the random reading, we could have to re do the same I/O later leading to a longer execution time. As explained previously, mapping the whole file will not bring it in RAM, only the parts of the file we want to access will be brought. Moreover, we have the possibility to "re use" previous I/O's when we need to access to an already accessed (and mapped) region.

One could imagine the best case is when  $B$  is equal to the total number of characters in the file. In order to prove this statement, a graph is built by plotting the execution time of the `randjump4` function with a constant number of iterations  $j$  equal to 10000, in function of  $B$  values going from 1310720000 ( $20000 \times 65536$ ) to 2129920000 ( $32500 \times 65536$ ) with a step of 655360 (10 times 65536) on the `cast_info` file, the following graph is obtained :

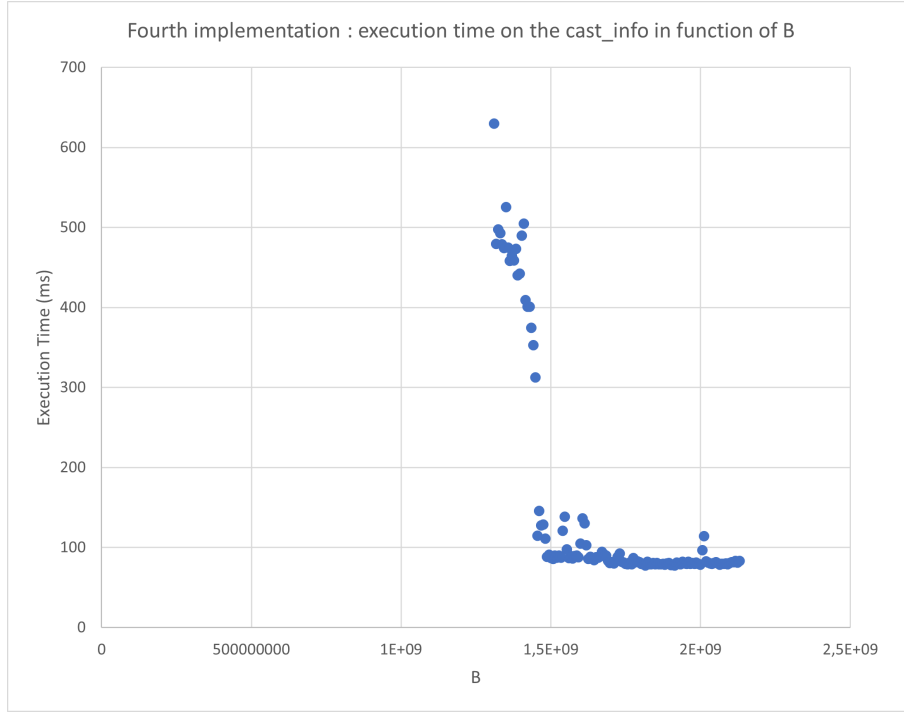


Figure 18: Execution time of the randjump function in Experiment 2 using InputStream4 on the cast\_info file with a constant  $j = 10000$ , in function of  $B$  varying from 1310720000 to 2129920000

The results are displayed in the table 37 in the Appendix B.4.

The graph seems constant after a sharp drop around  $B = 1448345600$  (22100 times 65536) where the execution time is equal to 312,412 ms. It is due to the fact that the total size of the `cast_info` file is equal to 1418137141, therefore, when  $B$  reaches approximately this value, the execution time drops. Thus, the execution time is drastically low in comparison with the previous results in the Figure 17. Concerning the execution times observed, they are smaller than the ones observed on the same file with the third implementation.

## 6.5 Comparison of the different implementations

After observing and collecting the data of the different implementations, it is obvious that the random reading of a file using the first implementation of `InputStream` is definitely the least efficient as for the first experiment.

By considering only the execution time of the 4 implementations with small  $B$  values, the third implementation is the most efficient. Indeed, the third implementation with a small size  $B$  will be the fastest among the 4 implementations. A comparative table is built below to justify this statement. The results are taken from the measurements on the `cast_info` file with a number of iterations  $j$  equal to 10000 and a seed equal 10, note that the size of the buffer  $B$  is equal to 10 for the third implementation and equal to 65536 for the fourth implementation.

| Implementation   | Execution time in ms (t) |
|------------------|--------------------------|
| implementation 1 | 797.049                  |
| implementation 2 | 248.819                  |
| implementation 3 | 142.298                  |
| implementation 4 | 544.931                  |

Table 15: Execution time measured from randjump function using the 4 `InputStream` on the `cast_info` file, with  $j = 10000$ ,  $B = 10$  (third implementation) and  $B = 65536$  (fourth implementation) and a seed = 10

The fastest implementation is indeed the third one as shown on the table 15.

However, by taking into account bigger size  $B$ , the result is different. Indeed, the fourth implementation is the most efficient when  $B$  is very big, as shown on the table 18 we can see when  $j$  is equal to 10000, the execution time is extremely faster in comparison to the others implementations especially the third one. To illustrate the comparison, the table 16 is built. The execution times are taken from the measurements on the `cast_info` file of the 4 implementations with the same parameters than the previous table except for the size of the buffer  $B$  of the fourth implementation which will be equal to 1500774400 (22900 times 65536).

| Implementation   | Execution time in ms (t) |
|------------------|--------------------------|
| implementation 1 | 797.049                  |
| implementation 2 | 248.819                  |
| implementation 3 | 142.298                  |
| implementation 4 | 86.8736                  |

Table 16: Execution time measured from randjump function using the 4 `InputStream` on the `cast_info` file, with  $j = 10000$ ,  $B = 10$  (third implementation) and  $B = 1500774400$  (fourth implementation) and a seed = 10

The fourth implementation is nearly 2 times faster than the third implementation. This result was expected, as the mapping is particular efficient for random reading.

## 7 Experiment 3 : Combined reading and writing

In this experiment, we are combining  $k$  different files into one unique file. To do so, we read one line of each input file and write them in an output file in a round robin fashion. The goal is to determine which pair of input stream and output stream works the best. For the input streams, we only considered the second and fourth implementation (respectively the best implementation identified in experiment 1 and 2).

All the analysis in the experiment 1 of the different implementation of `InputStream` can be transposed for the four implementations of `OutputStream`.

Before starting this experiment, we thought that `InputStream2` was going to be the fastest way to read the files because of the sequential reading. Indeed, even if we read the files in a round robin fashion, each file is still read sequentially. Therefore, the mapping would not bring any added value.

For the fastest way to write in the files, we supposed that it was going to be `OutputStream4` or `OutputStream2` based on the study made in the first experiment. Theoretically, the third implementation of `OutputStream` with a enormous `B` would have been the fastest one, but as we saw, the increase of speed with `B` will normally stop at a certain point. The first implementation of `OutputStream` will for sure be the slowest (based on what we studied in the first experiment). The advantage of the fourth implementation is that an I/O occurs only when we access a part of the file we want to write, as explained in section 4. Then, we modify in the RAM the buffer by copying our different lines without doing any I/O's. All the I/O's occurs only when unmaping or after a long time to update the content of the disk. It allows to have a faster way to write but, because we do less I/O's, if a crash occurs, it is more likely that we lose what we have modified.

Based on what we saw, the second implementation of `OutputStream` can be the fastest one, but the cache can only store a limited amount of bytes before flushing everything to the disk. The advantage of the cache when writing is less important than when reading as the OS cannot bring data while working on other process as in experiment 1 when reading.

However, if the cache can store and wait long enough to limit the number of I/O's, it can be nearly as fast as the fourth `OutputStream` implementation. To wrap up everything, we thought that because the cache does not add as much value as when reading, the fourth implementation of `OutputStream` was going to be the fastest way to write in the file.

Here, implementations that have a buffer are even more memory consuming (RAM memory) when using a big `k`. Indeed, we do not read a whole file and then read another one, instead we keep the buffer of every file in RAM and proceed to the reading in a round robin fashion. Because the mapping works with multiples of 16 pages, it can also become quickly memory consuming.

The first implementation of output stream is obviously less memory (RAM) consuming but also a lot slower because we do an I/O for each line read.

Different sets of files were created in order to measure the times of execution of this experiment, each letter corresponding to a set as displayed below. The further in the alphabetical order the letter is, the bigger sum of the file sizes is.

A = link\_type,kind\_type

B = link\_type,kind\_type,info\_type,company\_type,comp\_cast\_type,role\_type,movie\_link

C = complete\_cast,keyword

D = keyword,company\_name,movie\_info\_idx,movie\_keyword,movie\_companies,aka\_title,aka\_name

E = complete\_cast,keyword,movie\_link,company\_name,movie\_info\_idx,link\_type,kind\_type,info\_type,company\_type,comp\_cast\_type,role\_type,movie\_keyword,movie\_companies,aka\_title,aka\_name

F = movie\_info,person\_info

G = movie\_info,person\_info,name,cast\_info,title,movie\_keyword,char\_name

H = title,link\_type,kind\_type,info\_type,company\_type,comp\_cast\_type,role\_type,complete\_cast,keyword,movie\_link,company\_name,movie\_info\_idx,movie\_keyword,movie\_companies,

aka\_title,aka\_name,movie\_info,person\_info,  
name,cast\_info,char\_name

The parameter  $k$  correspond the number of files.

- $k=2$  for A,C and F.
- $k=7$  for B,D and G.
- $k=15$  for E.
- $k=21$  for H.

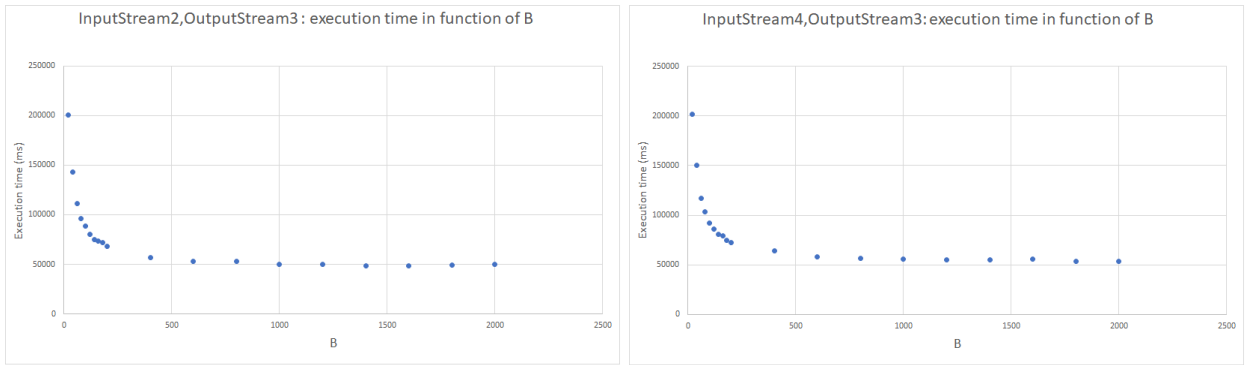
## 7.1 Variation of B

As we can observe in the graphs below, we have result similar to the ones we got in experiment 1 for `InputStream3` and `InputStream4`. The analysis done in experiment 1 for `InputStream3` and `InputStream4` can be transposed correspondingly for `OutputStream3` and `OutputStream4`.

The execution time drops as we increase B for `OutputStream3` for both pairs `InputStream2/OutputStream3` and `InputStream4/OutputStream3` until a stagnation at around  $B=1000$ . However, the execution time stays constant as we increase B for `OutputStream4` for both pairs `InputStream2/OutputStream4` and `InputStream4/OutputStream4`.

Note that the use of `InputStream2` or `InputStream4` does not influence the shape of the graph when changing B of the `OutputStream` peer (as expected).


In the next analysis, we will consider the execution time with  $B=2000$  for every pair containing `OutputStream3` and the mean execution time of all the B values for every pair containing `OutputStream4`.



(a) `InputStream2,OutputStream3`

(b) `InputStream4,OutputStream3`

Figure 19: Execution time in function of B on the set of files E



(a) InputStream2,OutputStream4

(b) InputStream4,OutputStream4

Figure 20: Execution time in function of B on the set of files E

## 7.2 Comparison of the pairs of InputStream/OutputStream

The execution times of this experiment on the different pairs of input and output streams, and applied to the different batches of files presented above are displayed :

| Implementation\Files       | A      | B        | C       | D       | E        | F        | G        | H        |
|----------------------------|--------|----------|---------|---------|----------|----------|----------|----------|
| InputStream2,OutputStream1 | 13.458 | 6114.02  |         |         |          |          |          |          |
| InputStream4,OutputStream1 | 13.909 | 6314.68  |         |         |          |          |          |          |
| InputStream2,OutputStream2 | 2.5805 | 116.485  | 1105.13 | 68432   | 50550.9  | 140699   | 444650   | 469675   |
| InputStream4,OutputStream2 | 2.7017 | 137.185  | 1215.01 | 89475.1 | 57618.2  | 144932   | 697026   | 706880   |
| InputStream2,OutputStream3 | 1.8721 | 113.727  | 1046.87 | 47096.9 | 49823.3  | 127123   | 427604   | 436578   |
| InputStream4,OutputStream3 | 2.0466 | 131.698  | 1198.47 | 54977   | 53551.9  | 136303   | 614177   | 623149   |
| InputStream2,OutputStream4 | 2.6061 | 113.4295 | 1049.21 | 42646.2 | 45608.6  | 110564.8 | 390415.2 | 4064633  |
| InputStream4,OutputStream4 | 2.9045 | 117.9432 | 1144.02 | 48628.8 | 48786.55 | 119011   | 572653.7 | 579699.2 |

Table 17: Time (in ms) in function of the files and the implementation

Remainder : for the third and fourth implementation of output streams, we respectively used in the table  $B = 2000$  and the average time on all the different B values that were tested (see appendix C).

Every empty cell corresponds to measurements that were not done because of the excessive amount of time it would have taken.

By analysing this table, we can see that the second implementation of input streams is always the fastest way to read. We can also see that for every set of files, the pair `InputStream2/OutputStream4` is always the fastest except for the column "A". Indeed, for very small files, the third implementation of `OutputStream` is the fastest way to write in the destination file. In fact this third implementation of `OutputStream` has almost the same execution time as the `OutputStream4` for columns B and C. We can deduce that by using a buffer size chosen by ourselves for `OutputStream3`, we can have similar or better execution time than `OutputStream4` which is not very optimal for small file sizes. However, when we work on very large files, `OutputStream4` becomes much more interesting. We can also see that `OutputStream2` does not perform as well as we expected, probably due to the fact that the cache can't optimize the writing as well as it does for the reading. As expected, `OutputStream1` is the slowest.

Also as expected, the execution time grows for every pairs of `InputStream` and `OutputStream` as the total size of the set grows (either by using a bigger k or by using bigger files).

## 8 Multi-way merge

In order to sort very large relations, it is possible to use the Two-Phase Multiway Merge-Sort algorithm. Considering the memory contains  $M$  buffers, the algorithm is separated into 2 phases :

- The file is divided into portions of  $M$  bytes, each portion of  $M$  bytes is written to a separate file. Notice that the lines in the file can't be separated, therefore lines are added to the file until its size exceeds  $M$ . These  $\left\lceil \frac{N}{M} \right\rceil$  files are sorted on the column  $k$ .
- $d$  files out of the  $\left\lceil \frac{N}{M} \right\rceil$  files are merged together in sorted order (following the merge-sort algorithm) with the help of a priority queue. This process is repeated until there is only one big sorted file left in the queue.


In order to evaluate the cost of this algorithm, we need to take into consideration the file size  $N$ , as well as  $M$  and  $d$ . The reading and writing functions used will be the optimal ones found in Experiment 3, which means the input streams used will be from the class `InputStream2` and the output streams will be from `OutputStream4`.

First of all, during the first phase, the whole file is read and written once. The cost of this phase is therefore the cost of reading  $N$  characters and writing  $N$  characters. This is therefore :

$$\left\lceil \frac{N}{S} \right\rceil + \left\lceil \frac{N}{G} \right\rceil$$

Indeed, the cost of reading a file sequentially using `InputStream2` was detailed in the Experiment 1, where  $S$  is the size of the internal buffer used by `fgets`, and the cost of writing  $N$  elements sequentially to a file using `OutputStream4` is the same as reading sequentially this file.

Next, these  $\left\lceil \frac{N}{M} \right\rceil$  files are merged together. In order to make the cost function simpler, we will provide an upper bound to the cost function, using the *big O* notation. Indeed, the problem encountered when figuring out the cost function is that if  $\left\lceil \frac{N}{M} \right\rceil \bmod d \neq 0$ , then the last files are merged with files that have already been merged previously (and are therefore longer). The cost function becomes in that case quite complicated to evaluate, especially that even when  $\left\lceil \frac{N}{M} \right\rceil \bmod d = 0$  it is very plausible that after a round or a few rounds we may end up with a number of files that is not a multiple of  $d$ . The *worst case scenario* is therefore when  $\left\lceil \frac{N}{M} \right\rceil \bmod d^r = 0$ , where  $r$  is the number of recursive *rounds* of merging that will be necessary. This is the case we will use in order to establish the cost function. Notice that, by focusing on this scenario, we can visualize the merging procedure as a tree. For clarity, let us represent the case where  $\left\lceil \frac{N}{M} \right\rceil = 8$  and  $d = 2$  :



In this tree,  $f_{i...j}$  corresponds to the file obtained by merging all files ranging from  $i$  to  $j$ . We can see that in total, the whole file has to be read and written 3 times in this example case. This helps us understand how to derive a function corresponding to the number of *rounds*  $r$  that are necessary to run the algorithm. Since the final result is always a single file, we can write :



$$\frac{\left\lceil \frac{N}{M} \right\rceil}{d^r} = 1 \quad (6)$$

from which we can establish :

$$r = \frac{\log \left\lceil \frac{N}{M} \right\rceil}{\log d} \quad (7)$$

which gives us the number of iterations of merging. Therefore, we can deduce the cost function of the multi-way merge sort algorithm we have implemented :

$$C(merge) = O \left( \frac{\log \left\lceil \frac{N}{M} \right\rceil}{\log d} \left( \left\lceil \frac{N}{S} \right\rceil + \left\lceil \frac{N}{G} \right\rceil \right) \right) \quad (8)$$

In order to verify the cost function, we will study independently the effects of variations of the parameters  $M$ ,  $N$  and  $d$ .

First, the execution time of the `extsort` function is called on the `company_name` file ( $N = 17802021$ ), with a constant  $d = 20$  and  $M$  varying between 10000 and 300000. Looking at the cost function, we would expect the execution time to decrease when  $M$  increases, following an inverse logarithmic trend, because of the factor  $\log \left\lceil \frac{N}{M} \right\rceil$  which is the only factor containing  $M$  in the cost formula. The main findings are showed in the table below (full results at appendix D) :

| $M$    | Execution time (t) | $t/\log(\text{ceil}(N/M))$ |
|--------|--------------------|----------------------------|
| 40000  | 145973             | 55098                      |
| 70000  | 134600             | 55931                      |
| 110000 | 129970             | 58822                      |
| 190000 | 113753             | 57651                      |

Table 18: Execution time of the multi-way merge sort algorithm on `company_name` in function of  $M$

The product  $\frac{t}{\log \left\lceil \frac{N}{M} \right\rceil}$  is computed at it should be a constant. By computing the average value as well as the standard deviation and coefficient of variation of this ratio for all computed values, we obtain :

$$\begin{aligned} \mu &= 57538,96 \\ \sigma &= 1880,17 \\ c_v &= \frac{\sigma}{\mu} = 3,27\% \end{aligned}$$

The coefficient of variation being extremely small, we can consider that our assumptions were correct. In order to further justify the dependency on  $M$  of the cost function, the results are plotted, with a logarithmic scale being used on the horizontal axis :




Figure 21: Execution time of the extsort function on the company\_name file in function of M, with constant values of d and N

We observe that the logarithmic trendline approximates the results quite accurately, and that its equation actually corresponds to an inverse logarithmic function. Therefore, we can consider that the execution time depends on M in an inverse logarithmic fashion as expected.

Next, the influence of N on the execution time has to be shown. Since this parameter appears in both factors of the cost function established at equation 8, we expect a  $N \cdot \log(N)$  evolution regarding the execution time when M and d are constant. Therefore, the multi-way merge sort algorithm was applied to most the files on the dataset of different lengths (the excessively long files were ignored due to the long execution times), while maintaining the other parameters constant. The very short files were also ignored as the M and d constant values (respectively 200000 and 200) did not really make sense for extremely short files. The main findings are shown below, and full results are available at appendix D :

| File name     | File length (N) | Execution time (t) | $N \cdot \log(N)$ | $N \cdot \log(\text{ceil}(N/M))/t$ |
|---------------|-----------------|--------------------|-------------------|------------------------------------|
| aka_name      | 73004383        | 694323             | 574059004         | 269,54                             |
| keyword       | 3791536         | 17931,8            | 24943815          | 270,4                              |
| movie_keyword | 93799307        | 1010650            | 747786800         | 247,91                             |
| title         | 203877939       | 2046890            | 1694097278        | 299,67                             |

Table 19: Execution time of the multi-way merge sort algorithm on several files in function of N, with constant values of M and d

The ratio  $\frac{N \cdot \log \left[ \frac{N}{M} \right]}{t}$  should behave more or less as a constant if we follow the cost function, although this is not entirely true as the N can not exactly be factored out of the second factor of equation 8. It is nevertheless a decent indicator to confirm that the established cost function is correct. The same statistical features as previously are computed :

$$\begin{aligned}\mu &= 237,58 \\ \sigma &= 63,88 \\ c_v &= \frac{\sigma}{\mu} = 26,8\%\end{aligned}$$

As expected, the coefficient of variation is higher than previously as this ratio does not exactly reflect the cost function, but it is still sufficiently low to confirm our hypothesis on the influence of  $N$  on the execution time. A graph plotting these execution times in function of  $N.\log(N)$  is presented :




Figure 22: Execution time of the extsort function on several files in function of  $N.\log(N)$ , with constant values of  $d$  and  $M$

The linear trendline seems to approximate the results well, and the intercept of this trendline is minimal (1396,1), therefore our assumption seems correct : the execution time of the multi-way merge sort evolves in function of  $N.\log(N)$ .

Finally, it is necessary to prove that the execution time of the multi-way merge sort algorithm depends on  $d$ , in an inverse logarithm fashion (see cost function at equation 8). Therefore, the execution times have been measured for  $d$  varying between 2 and 20 for the `company_name` file with  $M = 100000$ . The main results are shown below (full results at appendix D) :

| $d$ | Execution time (t) |
|-----|--------------------|
| 2   | 167956             |
| 5   | 139705             |
| 10  | 116257             |
| 14  | 111429             |

Table 20: Execution time of the multi-way merge sort algorithm on `company_name` in function of  $M$

By plotting the results of this execution in function of  $d$  on a graph with a logarithmic horizontal scale, we obtain :




Figure 23: Execution time of the extsort function on several files in function of  $d$ , with constant values of  $N$  and  $M$

We observe that the trendline seems to approximate the results pretty well, and that it indeed corresponds to an inverse logarithmic function. Therefore, we can assume that the multi-way merge sort execution time actually depends on  $d$  in an inverse logarithmic fashion as expected.

Therefore, we can conclude that the cost function established at equation 8 seems verified by the experimental results.

## 9 Conclusion

In this project, 4 different implementations for reading from and writing to a secondary memory have been done. By comparing their performance through experiments. After observing the results obtained from the first and second experiments, the two most efficient implementations for the reading of files were determined to be the second and the fourth implementations. This determination allowed to proceed with the next experiment which combined the reading implementations with the writing ones to establish the best pair of input stream implementation and output stream implementation. As a result, the second implementation of the input stream is the fastest way to read while for the output stream, the fourth implementation is more efficient with very large files and the third implementation is more interesting with small files. Finally, in order to fulfill the objectives of this project, an external-memory Two-Phase Multiway merge-sort algorithm has been implemented. Through different comparisons with plots, this merge-sort algorithm is confirmed to be dependent to the 3 parameters  $n$ ,  $d$  and  $M$ . Through this project the group has put into practice different concepts seen in the theoretical courses which has contributed a lot to our understanding.

## A Appendix 1 : Experiment 1

### A.1 First implementation

| File            | Average time (ms) | Length     | (Length/Time)      |
|-----------------|-------------------|------------|--------------------|
| aka_name        | 288342,2          | 73004383   | 253,1866061        |
| aka_title       | 155003            | 38858446   | 250,6947995        |
| cast_info       | 6804791           | 1418137141 | 208,4027476        |
| char_name       | 868975            | 215711567  | 248,2367928        |
| comp_cast_type  | 1,92637           | 45         | 23,35999834        |
| company_name    | 69193,9           | 17802021   | 257,2773178        |
| company_type    | 2,02237           | 92         | 45,49118114        |
| complete_cast   | 9242,43           | 2414495    | 261,2402799        |
| info_type       | 39,7328           | 1928       | 48,52414126        |
| keyword         | 13847,7           | 3791536    | 273,8025809        |
| kind_type       | 7,62033           | 85         | 11,15437258        |
| link_type       | 23,898            | 261        | 10,92141602        |
| movie_companies | 370852            | 93112104   | 251,0761813        |
| movie_info      | 3752617,4         | 792010495  | 211,0554876        |
| movie_info_idx  | 137217            | 35335875   | 257,5182011        |
| movie_keyword   | 362712            | 93799307   | 258,6054694        |
| movie_link      | 2611,33           | 656584     | 251,4366242        |
| name            | 1393964           | 321191609  | 230,416            |
| person_info     | 1595901,2         | 399133124  | 250,0988933        |
| role_type       | 4,12577           | 160        | 38,78063974        |
| title           | 856427            | 203877939  | 238,0564123        |
|                 |                   | <b>AVG</b> | <b>246,7402929</b> |
|                 |                   | <b>STD</b> | <b>17,9428158</b>  |
|                 |                   | <b>CV</b>  | <b>0,072719439</b> |

Figure 24: Execution time of the length function in Experiment 1 using the first implementation on all the files of the dataset

## A.2 Second implementation

| File            | Average time (ms) | Length     | (Length/Time)             |
|-----------------|-------------------|------------|---------------------------|
| aka_name        | 3667,61           | 73004383   | 19905,16522               |
| aka_title       | 1795,06           | 38858446   | 21647,43574               |
| cast_info       | 102248            | 1,418E+09  | 13869,58318               |
| char_name       | 11619,2           | 215711567  | 18565,09631               |
| comp_cast_type  | 0,151             | 45         | 298,013245                |
| company_name    | 921,501           | 17802021   | 19318,50427               |
| company_type    | 0,15554           | 92         | 591,4877202               |
| complete_cast   | 282,008           | 2414495    | 8561,796119               |
| info_type       | 0,35546           | 1928       | 5423,957689               |
| keyword         | 330,201           | 3791536    | 11482,50914               |
| kind_type       | 9,6749            | 85         | 8,785620523               |
| link_type       | 0,17404           | 261        | 1499,655252               |
| movie_companies | 6938,79           | 93112104   | 13419,06932               |
| movie_info      | 48641,7           | 792010495  | 16282,54142               |
| movie_info_idx  | 3240,73           | 35335875   | 10903,67757               |
| movie_keyword   | 9915,18           | 93799307   | 9460,171878               |
| movie_link      | 69,1148           | 656584     | 9499,904507               |
| name            | 16333,1           | 321191609  | 19665,07332               |
| person_info     | 28511,8           | 399133124  | 13998,87499               |
| role_type       | 0,17812           | 160        | 898,2708287               |
| title           | 10162,5           | 203877939  | 20061,78982               |
|                 |                   | <b>AVG</b> | <b><u>15109,41285</u></b> |
|                 |                   | <b>STD</b> | <b><u>4357,228925</u></b> |
|                 |                   | <b>CV</b>  | <b><u>0,288378441</u></b> |

Figure 25: Execution time of the length function in Experiment 1 using the second implementation on all the files of the dataset

### A.3 Third implementation

| File            | Average time (ms) | Length    | ceil(N/B)  | ceil(N/B)/t        |
|-----------------|-------------------|-----------|------------|--------------------|
| aka_name        | 8785,92           | 73004383  | 7300439    | 830,9248206        |
| aka_title       | 4530,44           | 38858446  | 3885845    | 857,7191178        |
| cast_info       | 189036            | 1,418E+09 | 141813715  | 750,1942223        |
| char_name       | 25043,9           | 215711567 | 21571157   | 861,3337779        |
| comp_cast_type  | 14,5047           | 45        | 5          | 0,344715851        |
| company_name    | 2034,13           | 17802021  | 1780203    | 875,1667789        |
| company_type    | 11,025            | 92        | 10         | 0,907029478        |
| complete_cast   | 338,033           | 2414495   | 241450     | 714,2793751        |
| info_type       | 18,7994           | 1928      | 193        | 10,26628509        |
| keyword         | 550,009           | 3791536   | 379154     | 689,3596287        |
| kind_type       | 0,95146           | 85        | 9          | 9,459146995        |
| link_type       | 4,581             | 261       | 27         | 5,893909627        |
| movie_companies | 12442,1           | 93112104  | 9311211    | 748,3632988        |
| movie_info      | 83438,9           | 792010495 | 79201050   | 949,2101406        |
| movie_info_idx  | 5144,85           | 35335875  | 3533588    | 686,8204126        |
| movie_keyword   | 15221,9           | 93799307  | 9379931    | 616,2128906        |
| movie_link      | 82,332            | 656584    | 65659      | 797,4906476        |
| name            | 37347,6           | 321191609 | 32119161   | 860,0060245        |
| person_info     | 37769,1           | 399133124 | 39913313   | 1056,77162         |
| role_type       | 10,5923           | 160       | 16         | 1,510531235        |
| title           | 23236,6           | 203877939 | 20387794   | 877,4000499        |
|                 |                   |           | <b>AVG</b> | <b>811,4168537</b> |
|                 |                   |           | <b>STD</b> | <b>109,8062751</b> |
|                 |                   |           | <b>CV</b>  | <b>0,135326589</b> |

Figure 26: Execution time of the length function in Experiment 1 using the third implementation with B=100 on all the files of the dataset

| B    | t       | ceil(N/B) | ceil(N/B)/t |
|------|---------|-----------|-------------|
| 10   | 31596,1 | 7300439   | 231,0550669 |
| 20   | 19525,9 | 3650220   | 186,9424713 |
| 30   | 14372,8 | 2433480   | 169,3114772 |
| 40   | 12207,9 | 1825110   | 149,5023714 |
| 50   | 10836,9 | 1460088   | 134,7329956 |
| 60   | 10039,9 | 1216740   | 121,1904501 |
| 70   | 9355,44 | 1042920   | 111,4773864 |
| 80   | 8931,16 | 912555    | 102,176537  |
| 90   | 8484,24 | 811160    | 95,60785645 |
| 100  | 8151,63 | 730044    | 89,55803931 |
| 110  | 7906,95 | 663677    | 83,93590449 |
| 120  | 8130,3  | 608370    | 74,82749714 |
| 130  | 7507,11 | 561573    | 74,8054844  |
| 140  | 7320,34 | 521460    | 71,23439622 |
| 150  | 7225,41 | 486696    | 67,35894572 |
| 160  | 7054,26 | 456278    | 64,68119973 |
| 170  | 6962,26 | 429438    | 61,68083352 |
| 180  | 6834,09 | 405580    | 59,34659918 |
| 190  | 6760,15 | 384234    | 56,83808791 |
| 200  | 6692,75 | 365022    | 54,53991259 |
| 210  | 6612,3  | 347640    | 52,57474706 |
| 220  | 6566,44 | 331839    | 50,53560224 |
| 230  | 6518,69 | 317411    | 48,69245201 |
| 240  | 6469,82 | 304185    | 47,01599117 |
| 250  | 6405,8  | 292018    | 45,58649973 |
| 260  | 6335,56 | 280787    | 44,31920777 |
| 270  | 6773,49 | 270387    | 39,91841724 |
| 280  | 6271,93 | 260730    | 41,5709359  |
| 290  | 6220,33 | 251740    | 40,47052166 |
| 300  | 6221,11 | 243348    | 39,11649207 |
| 310  | 6416,13 | 235499    | 36,70421266 |
| 320  | 6133,76 | 228139    | 37,19398868 |
| 330  | 6168,23 | 221226    | 35,86539412 |
| 340  | 6087,27 | 214719    | 35,2734477  |
| 350  | 6047,58 | 208584    | 34,49049041 |
| 360  | 6126    | 202790    | 33,10316683 |
| 370  | 6072,04 | 197310    | 32,49484522 |
| 380  | 5980,69 | 192117    | 32,12288214 |
| 390  | 5975,13 | 187191    | 31,32835604 |
| 400  | 5938,87 | 182511    | 30,73160382 |
| 410  | 5915,56 | 178060    | 30,10027791 |
| 420  | 5932,06 | 173820    | 29,30179398 |
| 430  | 5917,52 | 169778    | 28,69073531 |
| 440  | 6363,75 | 165920    | 26,07267727 |
| 450  | 5912,75 | 162232    | 27,43765591 |
| 460  | 5841,5  | 158706    | 27,16870667 |
| 470  | 5904,99 | 155329    | 26,30470162 |
| 480  | 5823,7  | 152093    | 26,11621478 |
| 490  | 5814,98 | 148989    | 25,62158425 |
| 500  | 5877,3  | 146009    | 24,84287003 |
| 510  | 5778,5  | 143146    | 24,77217271 |
| 520  | 5790,4  | 140394    | 24,24599337 |
| 530  | 5788,85 | 137745    | 23,79488154 |
| 540  | 5780,01 | 135194    | 23,38992493 |
| 550  | 5761,76 | 132736    | 23,03740524 |
| 560  | 5748,25 | 130365    | 22,67907624 |
| 570  | 5769,25 | 128078    | 22,20011267 |
| 580  | 5715,02 | 125870    | 22,02441986 |
| 590  | 5706,23 | 123737    | 21,68454479 |
| 600  | 5720,09 | 121674    | 21,27134363 |
| 610  | 5716,85 | 119680    | 20,9346056  |
| 620  | 6147,33 | 117750    | 19,15465739 |
| 630  | 5685,1  | 115880    | 20,38310672 |
| 640  | 5739,34 | 114070    | 19,87510759 |
| 650  | 5678,24 | 112315    | 19,77989659 |
| 660  | 5691,18 | 110613    | 19,43586392 |
| 670  | 6067,21 | 108962    | 17,9591608  |
| 680  | 5651,24 | 107360    | 18,99760053 |
| 690  | 5635,88 | 105804    | 18,77328829 |
| 700  | 5619,73 | 104292    | 18,55818696 |
| 710  | 5708,56 | 102824    | 18,01224827 |
| 720  | 5615,39 | 101395    | 18,05662652 |
| 730  | 5636,69 | 100007    | 17,74215009 |
| 740  | 5691,82 | 98655     | 17,33276878 |
| 750  | 5648,98 | 97340     | 17,23142939 |
| 760  | 5614,12 | 96059     | 17,11025058 |
| 770  | 5605,87 | 94811     | 16,91280747 |
| 780  | 5610,79 | 93596     | 16,68142989 |
| 790  | 5605,55 | 92411     | 16,48562585 |
| 800  | 6115,44 | 91256     | 14,92222964 |
| 810  | 5648,15 | 90129     | 15,95726034 |
| 820  | 5570,03 | 89030     | 15,98375592 |
| 830  | 5887,21 | 87958     | 14,94052361 |
| 840  | 5569,51 | 86910     | 15,60460435 |
| 850  | 5595,2  | 85888     | 15,35030026 |
| 860  | 5590,05 | 84889     | 15,18573179 |
| 870  | 5557,53 | 83914     | 15,09915376 |
| 880  | 5645,39 | 82960     | 14,69517606 |
| 890  | 5603,96 | 82028     | 14,63750633 |
| 900  | 5558,77 | 81116     | 14,59243682 |
| 910  | 5608,13 | 80225     | 14,30512488 |
| 920  | 5550,57 | 79353     | 14,29636956 |
| 930  | 5534,9  | 78500     | 14,1827314  |
| 940  | 5557,77 | 77665     | 13,97412991 |
| 950  | 5546,97 | 76847     | 13,85386977 |
| 960  | 5924,26 | 76047     | 12,83653992 |
| 970  | 5551,78 | 75263     | 13,55655303 |
| 980  | 6045,61 | 74495     | 12,32216435 |
| 990  | 5587,31 | 73742     | 13,19812217 |
| 1000 | 5566,24 | 73005     | 13,11567593 |

Figure 27: Execution time of the length function in Experiment 1 using the third implementation on `aka_name` with `B` varying from 10 to 1000



#### A.4 Fourth implementation

| File            | Average time (ms) | Length     | ceil(N/65536) | ceil(N/65536)/t    |
|-----------------|-------------------|------------|---------------|--------------------|
| aka_name        | 5361,78           | 73004383   | 1114          | 0,207766824        |
| aka_title       | 2521,57           | 38858446   | 593           | 0,235170945        |
| cast_info       | 127241            | 1418137141 | 21640         | 0,170070968        |
| char_name       | 16252,8           | 215711567  | 3292          | 0,202549715        |
| comp_cast_type  | 0,25072           | 45         | 1             | 3,988513082        |
| company_name    | 1290,66           | 17802021   | 272           | 0,21074489         |
| company_type    | 0,25496           | 92         | 1             | 3,922183872        |
| complete_cast   | 317,221           | 2414495    | 37            | 0,116637928        |
| info_type       | 10,9712           | 1928       | 1             | 0,091147732        |
| keyword         | 369,457           | 3791536    | 58            | 0,156987146        |
| kind_type       | 0,25826           | 85         | 1             | 3,872066909        |
| link_type       | 0,27686           | 261        | 1             | 3,611933829        |
| movie_companies | 8466,79           | 93112104   | 1421          | 0,167832201        |
| movie_info      | 51280,3           | 792010495  | 12086         | 0,235685049        |
| movie_info_idx  | 3511,95           | 35335875   | 540           | 0,153760731        |
| movie_keyword   | 10786,7           | 93799307   | 1432          | 0,132756079        |
| movie_link      | 85,5874           | 656584     | 11            | 0,128523591        |
| name            | 23826,5           | 321191609  | 4901          | 0,205695339        |
| person_info     | 22285,3           | 399133124  | 6091          | 0,273319183        |
| role_type       | 0,2556            | 160        | 1             | 3,912363067        |
| title           | 14670,1           | 203877939  | 3111          | 0,212063994        |
|                 |                   |            | <b>AVG</b>    | <b>0,191502928</b> |
|                 |                   |            | <b>STD</b>    | <b>0,041939788</b> |
|                 |                   |            | <b>CV</b>     | <b>0,219003377</b> |

Figure 28: Execution time of the length function in Experiment 1 using the fourth implementation on all the files of the dataset

## A.5 Comparisons

| B    | LENGTH1  | LENGTH2 | LENGTH3 | LENGTH4 |
|------|----------|---------|---------|---------|
| 10   | 288342,2 | 3587,4  | 31596,1 | 5361,78 |
| 20   | 288342,2 | 3587,4  | 19525,9 | 5361,78 |
| 30   | 288342,2 | 3587,4  | 14372,8 | 5361,78 |
| 40   | 288342,2 | 3587,4  | 12207,9 | 5361,78 |
| 50   | 288342,2 | 3587,4  | 10836,9 | 5361,78 |
| 60   | 288342,2 | 3587,4  | 10039,9 | 5361,78 |
| 70   | 288342,2 | 3587,4  | 9355,44 | 5361,78 |
| 80   | 288342,2 | 3587,4  | 8931,16 | 5361,78 |
| 90   | 288342,2 | 3587,4  | 8484,24 | 5361,78 |
| 100  | 288342,2 | 3587,4  | 8151,63 | 5361,78 |
| 110  | 288342,2 | 3587,4  | 7906,95 | 5361,78 |
| 120  | 288342,2 | 3587,4  | 8130,3  | 5361,78 |
| 130  | 288342,2 | 3587,4  | 7507,11 | 5361,78 |
| 140  | 288342,2 | 3587,4  | 7320,34 | 5361,78 |
| 150  | 288342,2 | 3587,4  | 7225,41 | 5361,78 |
| 160  | 288342,2 | 3587,4  | 7054,26 | 5361,78 |
| 170  | 288342,2 | 3587,4  | 6962,26 | 5361,78 |
| 180  | 288342,2 | 3587,4  | 6834,09 | 5361,78 |
| 190  | 288342,2 | 3587,4  | 6760,15 | 5361,78 |
| 200  | 288342,2 | 3587,4  | 6692,75 | 5361,78 |
| 210  | 288342,2 | 3587,4  | 6612,3  | 5361,78 |
| 220  | 288342,2 | 3587,4  | 6566,44 | 5361,78 |
| 230  | 288342,2 | 3587,4  | 6518,69 | 5361,78 |
| 240  | 288342,2 | 3587,4  | 6469,82 | 5361,78 |
| 250  | 288342,2 | 3587,4  | 6405,8  | 5361,78 |
| 260  | 288342,2 | 3587,4  | 6335,56 | 5361,78 |
| 270  | 288342,2 | 3587,4  | 6773,49 | 5361,78 |
| 280  | 288342,2 | 3587,4  | 6271,93 | 5361,78 |
| 290  | 288342,2 | 3587,4  | 6220,33 | 5361,78 |
| 300  | 288342,2 | 3587,4  | 6221,11 | 5361,78 |
| 310  | 288342,2 | 3587,4  | 6416,13 | 5361,78 |
| 320  | 288342,2 | 3587,4  | 6133,76 | 5361,78 |
| 330  | 288342,2 | 3587,4  | 6168,23 | 5361,78 |
| 340  | 288342,2 | 3587,4  | 6087,27 | 5361,78 |
| 350  | 288342,2 | 3587,4  | 6047,58 | 5361,78 |
| 360  | 288342,2 | 3587,4  | 6126    | 5361,78 |
| 370  | 288342,2 | 3587,4  | 6072,04 | 5361,78 |
| 380  | 288342,2 | 3587,4  | 5980,69 | 5361,78 |
| 390  | 288342,2 | 3587,4  | 5975,13 | 5361,78 |
| 400  | 288342,2 | 3587,4  | 5938,87 | 5361,78 |
| 410  | 288342,2 | 3587,4  | 5915,56 | 5361,78 |
| 420  | 288342,2 | 3587,4  | 5932,06 | 5361,78 |
| 430  | 288342,2 | 3587,4  | 5917,52 | 5361,78 |
| 440  | 288342,2 | 3587,4  | 6363,75 | 5361,78 |
| 450  | 288342,2 | 3587,4  | 5912,75 | 5361,78 |
| 460  | 288342,2 | 3587,4  | 5841,5  | 5361,78 |
| 470  | 288342,2 | 3587,4  | 5904,99 | 5361,78 |
| 480  | 288342,2 | 3587,4  | 5823,7  | 5361,78 |
| 490  | 288342,2 | 3587,4  | 5814,98 | 5361,78 |
| 500  | 288342,2 | 3587,4  | 5877,3  | 5361,78 |
| 510  | 288342,2 | 3587,4  | 5778,5  | 5361,78 |
| 520  | 288342,2 | 3587,4  | 5790,4  | 5361,78 |
| 530  | 288342,2 | 3587,4  | 5788,85 | 5361,78 |
| 540  | 288342,2 | 3587,4  | 5780,01 | 5361,78 |
| 550  | 288342,2 | 3587,4  | 5761,76 | 5361,78 |
| 560  | 288342,2 | 3587,4  | 5748,25 | 5361,78 |
| 570  | 288342,2 | 3587,4  | 5769,25 | 5361,78 |
| 580  | 288342,2 | 3587,4  | 5715,02 | 5361,78 |
| 590  | 288342,2 | 3587,4  | 5706,23 | 5361,78 |
| 600  | 288342,2 | 3587,4  | 5720,09 | 5361,78 |
| 610  | 288342,2 | 3587,4  | 5716,85 | 5361,78 |
| 620  | 288342,2 | 3587,4  | 6147,33 | 5361,78 |
| 630  | 288342,2 | 3587,4  | 5685,1  | 5361,78 |
| 640  | 288342,2 | 3587,4  | 5739,34 | 5361,78 |
| 650  | 288342,2 | 3587,4  | 5678,24 | 5361,78 |
| 660  | 288342,2 | 3587,4  | 5691,18 | 5361,78 |
| 670  | 288342,2 | 3587,4  | 6067,21 | 5361,78 |
| 680  | 288342,2 | 3587,4  | 5651,24 | 5361,78 |
| 690  | 288342,2 | 3587,4  | 5635,88 | 5361,78 |
| 700  | 288342,2 | 3587,4  | 5619,73 | 5361,78 |
| 710  | 288342,2 | 3587,4  | 5708,56 | 5361,78 |
| 720  | 288342,2 | 3587,4  | 5615,39 | 5361,78 |
| 730  | 288342,2 | 3587,4  | 5636,69 | 5361,78 |
| 740  | 288342,2 | 3587,4  | 5691,82 | 5361,78 |
| 750  | 288342,2 | 3587,4  | 5648,98 | 5361,78 |
| 760  | 288342,2 | 3587,4  | 5614,12 | 5361,78 |
| 770  | 288342,2 | 3587,4  | 5605,87 | 5361,78 |
| 780  | 288342,2 | 3587,4  | 5610,79 | 5361,78 |
| 790  | 288342,2 | 3587,4  | 5605,55 | 5361,78 |
| 800  | 288342,2 | 3587,4  | 6115,44 | 5361,78 |
| 810  | 288342,2 | 3587,4  | 5648,15 | 5361,78 |
| 820  | 288342,2 | 3587,4  | 5570,03 | 5361,78 |
| 830  | 288342,2 | 3587,4  | 5887,21 | 5361,78 |
| 840  | 288342,2 | 3587,4  | 5569,51 | 5361,78 |
| 850  | 288342,2 | 3587,4  | 5595,2  | 5361,78 |
| 860  | 288342,2 | 3587,4  | 5590,05 | 5361,78 |
| 870  | 288342,2 | 3587,4  | 5557,53 | 5361,78 |
| 880  | 288342,2 | 3587,4  | 5645,39 | 5361,78 |
| 890  | 288342,2 | 3587,4  | 5603,96 | 5361,78 |
| 900  | 288342,2 | 3587,4  | 5558,77 | 5361,78 |
| 910  | 288342,2 | 3587,4  | 5608,13 | 5361,78 |
| 920  | 288342,2 | 3587,4  | 5550,57 | 5361,78 |
| 930  | 288342,2 | 3587,4  | 5534,9  | 5361,78 |
| 940  | 288342,2 | 3587,4  | 5557,77 | 5361,78 |
| 950  | 288342,2 | 3587,4  | 5546,97 | 5361,78 |
| 960  | 288342,2 | 3587,4  | 5924,26 | 5361,78 |
| 970  | 288342,2 | 3587,4  | 5551,78 | 5361,78 |
| 980  | 288342,2 | 3587,4  | 6045,61 | 5361,78 |
| 990  | 288342,2 | 3587,4  | 5587,31 | 5361,78 |
| 1000 | 288342,2 | 3587,4  | 5566,24 | 5361,78 |

Figure 29: Execution time of the length function in Experiment 1 using all the implementations on aka\_name with B varying from 10 to 1000

## B Appendix 2 : Experiment 2

### B.1 First implementation

| Iterations | Average time (ms) | Sum    |
|------------|-------------------|--------|
| 100        | 19,8876           | 2052   |
| 200        | 25,2662           | 4207   |
| 300        | 25,9037           | 6214   |
| 400        | 32,0613           | 8082   |
| 500        | 44,8529           | 10261  |
| 600        | 48,7866           | 12366  |
| 700        | 57,9411           | 14298  |
| 800        | 72,3914           | 16403  |
| 900        | 76,7728           | 18564  |
| 1000       | 81,9804           | 20585  |
| 2000       | 159,575           | 41328  |
| 3000       | 242,634           | 62234  |
| 4000       | 320,111           | 83300  |
| 5000       | 404,636           | 104084 |
| 6000       | 485,65            | 124624 |
| 7000       | 557,904           | 144875 |
| 8000       | 638,603           | 165627 |
| 9000       | 717,297           | 186091 |
| 10000      | 797,049           | 206576 |

Figure 30: Execution time of the randjump function in Experiment 2 using the first implementation on the cast\_info file in function of large number of iterations

## B.2 Second implementation

| Iterations | Average time (ms) | Sum    |
|------------|-------------------|--------|
| 100        | 4,77173           | 2052   |
| 200        | 18,1898           | 4207   |
| 300        | 16,0279           | 6214   |
| 400        | 13,8294           | 8082   |
| 500        | 18,8686           | 10261  |
| 600        | 21,3854           | 12366  |
| 700        | 21,3214           | 14298  |
| 800        | 25,7332           | 16403  |
| 900        | 22,7223           | 18564  |
| 1000       | 29,1971           | 20585  |
| 2000       | 143,523           | 41328  |
| 3000       | 153,204           | 62234  |
| 4000       | 173,845           | 83300  |
| 5000       | 197,101           | 104084 |
| 6000       | 187,592           | 124624 |
| 7000       | 207,32            | 144875 |
| 8000       | 264,937           | 165627 |
| 9000       | 192,719           | 186091 |
| 10000      | 248,819           | 206576 |

Figure 31: Execution time of the randjump function in Experiment 2 using the second implementation on the cast\_info file in function of large number of iterations

### B.3 Third implementation

| Iterations | Average time (ms) | Sum    |
|------------|-------------------|--------|
| 100        | 7,65031           | 2165   |
| 200        | 14,2564           | 4242   |
| 300        | 16,7692           | 6425   |
| 400        | 24,7671           | 8551   |
| 500        | 33,981            | 10636  |
| 600        | 44,3552           | 12978  |
| 700        | 41,499            | 15141  |
| 800        | 44,2708           | 17323  |
| 900        | 51,1397           | 19537  |
| 1000       | 56,4301           | 21679  |
| 2000       | 131,35            | 43639  |
| 3000       | 167,318           | 65077  |
| 4000       | 227,703           | 87102  |
| 5000       | 280,503           | 109122 |
| 6000       | 338,069           | 130937 |
| 7000       | 390,28            | 152509 |
| 8000       | 444,053           | 173438 |
| 9000       | 499,736           | 195409 |
| 10000      | 559,267           | 217366 |

Figure 32: Execution time of the randjump function in Experiment 2 using the third implementation on the cast\_info file with a constant  $B = 100$ , in function of large number of iterations

| B   | Average time (ms) | Sum    |      |          |        |
|-----|-------------------|--------|------|----------|--------|
| 10  | 142,298           | 206576 | 540  | 159,6746 | 206576 |
| 20  | 154,372           | 206576 | 550  | 144,3604 | 206576 |
| 30  | 144,53            | 206576 | 560  | 167,9912 | 206576 |
| 40  | 142,891           | 206576 | 570  | 168,5996 | 206576 |
| 50  | 137,066           | 206576 | 580  | 162,7522 | 206576 |
| 60  | 134,901           | 206576 | 590  | 148,689  | 206576 |
| 70  | 135,756           | 206576 | 600  | 149,8984 | 206576 |
| 80  | 135,21            | 206576 | 610  | 159,826  | 206576 |
| 90  | 135,738           | 206576 | 620  | 150,4144 | 206576 |
| 100 | 140,212           | 206576 | 630  | 142,7794 | 206576 |
| 110 | 137,7             | 206576 | 640  | 176,9776 | 206576 |
| 120 | 137,46            | 206576 | 650  | 160,4942 | 206576 |
| 130 | 138,374           | 206576 | 660  | 162,5026 | 206576 |
| 140 | 141,517           | 206576 | 670  | 178,8442 | 206576 |
| 150 | 138,597           | 206576 | 680  | 180,4184 | 206576 |
| 160 | 142,386           | 206576 | 690  | 172,2814 | 206576 |
| 170 | 145,293           | 206576 | 700  | 169,2222 | 206576 |
| 180 | 146,928           | 206576 | 710  | 174,0662 | 206576 |
| 190 | 159,02            | 206576 | 720  | 166,6324 | 206576 |
| 200 | 149,41            | 206576 | 730  | 205,104  | 206576 |
| 210 | 146,569           | 206576 | 740  | 172,6408 | 206576 |
| 220 | 146,819           | 206576 | 750  | 149,7412 | 206576 |
| 230 | 142,25            | 206576 | 760  | 161,6194 | 206576 |
| 240 | 145,444           | 206576 | 770  | 162,873  | 206576 |
| 250 | 141,994           | 206576 | 780  | 180,6742 | 206576 |
| 260 | 147,991           | 206576 | 790  | 168,6382 | 206576 |
| 270 | 141,655           | 206576 | 800  | 156,1466 | 206576 |
| 280 | 149,536           | 206576 | 810  | 149,1708 | 206576 |
| 290 | 144,841           | 206576 | 820  | 171,1766 | 206576 |
| 300 | 147,541           | 206576 | 830  | 176,0656 | 206576 |
| 310 | 148,383           | 206576 | 840  | 177,805  | 206576 |
| 320 | 147,341           | 206576 | 850  | 200,614  | 206576 |
| 330 | 149,505           | 206576 | 860  | 170,585  | 206576 |
| 340 | 145,579           | 206576 | 870  | 172,3044 | 206576 |
| 350 | 146,857           | 206576 | 880  | 200,384  | 206576 |
| 360 | 147,044           | 206576 | 890  | 188,321  | 206576 |
| 370 | 149,463           | 206576 | 900  | 176,7914 | 206576 |
| 380 | 149,564           | 206576 | 910  | 165,2132 | 206576 |
| 390 | 148,083           | 206576 | 920  | 202,726  | 206576 |
| 400 | 148,997           | 206576 | 930  | 151,6324 | 206576 |
| 410 | 154,4038          | 206576 | 940  | 193,3314 | 206576 |
| 420 | 144,0886          | 206576 | 950  | 159,7178 | 206576 |
| 430 | 140,9236          | 206576 | 960  | 194,1226 | 206576 |
| 440 | 143,5764          | 206576 | 970  | 176,5678 | 206576 |
| 450 | 157,73            | 206576 | 980  | 178,7506 | 206576 |
| 460 | 143,853           | 206576 | 990  | 208,634  | 206576 |
| 470 | 161,0208          | 206576 | 1000 | 218,512  | 206576 |
| 480 | 144,0728          | 206576 |      |          |        |
| 490 | 154,1616          | 206576 |      |          |        |
| 500 | 141,1706          | 206576 |      |          |        |
| 510 | 154,1706          | 206576 |      |          |        |
| 520 | 140,6316          | 206576 |      |          |        |
| 530 | 158,1628          | 206576 |      |          |        |

Figure 33: Execution time of the randjump function in Experiment 2 using the third implementation on the cast\_info file with B varying from 10 to 1000 and  $j = 10000$

| B     | Average time (ms) | Sum    |       |         |        |
|-------|-------------------|--------|-------|---------|--------|
| 10000 | 486,919           | 206576 |       |         |        |
| 11000 | 501,472           | 206576 | 32000 | 1175,42 | 206576 |
| 12000 | 539,704           | 206576 | 33000 | 1208    | 206576 |
| 13000 | 564,197           | 206576 | 34000 | 1241,01 | 206576 |
| 14000 | 599,334           | 206576 | 35000 | 1268,45 | 206576 |
| 15000 | 632               | 206576 | 36000 | 1304,78 | 206576 |
| 16000 | 662,792           | 206576 | 37000 | 1332,82 | 206576 |
| 17000 | 693,385           | 206576 | 38000 | 1381,18 | 206576 |
| 18000 | 722,628           | 206576 | 39000 | 1410,8  | 206576 |
| 19000 | 761,774           | 206576 | 40000 | 1436,65 | 206576 |
| 20000 | 787,248           | 206576 | 41000 | 1471,24 | 206576 |
| 21000 | 853,991           | 206576 | 42000 | 1501,86 | 206576 |
| 22000 | 860,888           | 206576 | 43000 | 1697,47 | 206576 |
| 23000 | 883,573           | 206576 | 44000 | 1759,56 | 206576 |
| 24000 | 920,877           | 206576 | 45000 | 1612,43 | 206576 |
| 25000 | 948,203           | 206576 | 46000 | 1651,62 | 206576 |
| 26000 | 980,073           | 206576 | 47000 | 1686,55 | 206576 |
| 27000 | 1012,04           | 206576 | 48000 | 1680,54 | 206576 |
| 28000 | 1047,47           | 206576 | 49000 | 1673,08 | 206576 |
| 29000 | 1083,08           | 206576 | 50000 | 1995,72 | 206576 |
| 30000 | 1125,06           | 206576 |       |         |        |
| 31000 | 1229,48           | 206576 |       |         |        |

Figure 34: Execution time of the randjump function in Experiment 2 using the third implementation on the cast\_info file with B varying from 10000 to 50000 and  $j = 10000$

#### B.4 Fourth implementation

| Iterations | Average ti | Sum    |
|------------|------------|--------|
| 100        | 8,35968    | 2150   |
| 200        | 14,0492    | 4401   |
| 300        | 18,5894    | 6504   |
| 400        | 28,2143    | 8469   |
| 500        | 30,2253    | 10743  |
| 600        | 36,9993    | 12942  |
| 700        | 49,182     | 14972  |
| 800        | 49,8882    | 17175  |
| 900        | 56,3039    | 19432  |
| 1000       | 58,7572    | 21551  |
| 2000       | 135,484    | 43271  |
| 3000       | 181,076    | 65154  |
| 4000       | 240,315    | 87199  |
| 5000       | 278,724    | 108963 |
| 6000       | 327,128    | 130471 |
| 7000       | 386,82     | 151698 |
| 8000       | 441,374    | 173433 |
| 9000       | 500,439    | 194871 |
| 10000      | 544,931    | 216325 |

Figure 35: Execution time of the randjump function in Experiment 2 using the fourth implementation on the cast\_info file with a constant  $B = 65536$ , in function of large number of iterations



| B        | Average time (ms) | Sum    |          |         |        |          |         |        |
|----------|-------------------|--------|----------|---------|--------|----------|---------|--------|
| 655360   | 574,446           | 207601 | 27525120 | 662,548 | 207601 | 54394880 | 615,715 | 207601 |
| 1310720  | 642,694           | 207601 | 28180480 | 636,47  | 207601 | 55050240 | 615,529 | 207601 |
| 1966080  | 696,489           | 207601 | 28835840 | 621,618 | 207601 | 55705600 | 614,66  | 207601 |
| 2621440  | 619,632           | 207601 | 29491200 | 617,671 | 207601 | 56360960 | 614,66  | 207601 |
| 3276800  | 615,8             | 207601 | 30146560 | 617,21  | 207601 | 57016320 | 615,116 | 207601 |
| 3932160  | 611,985           | 207601 | 30801920 | 617,922 | 207601 | 57671680 | 615,836 | 207601 |
| 4587520  | 618,978           | 207601 | 31457280 | 616,272 | 207601 | 58327040 | 615,45  | 207601 |
| 5242880  | 609,566           | 207601 | 32112640 | 619,544 | 207601 | 58982400 | 625,746 | 207601 |
| 5898240  | 614,761           | 207601 | 32768000 | 618,077 | 207601 | 59637760 | 655,996 | 207601 |
| 6553600  | 616,963           | 207601 | 33423360 | 623,435 | 207601 | 60293120 | 652,287 | 207601 |
| 7208960  | 621,637           | 207601 | 34078720 | 618,636 | 207601 | 60948480 | 615,475 | 207601 |
| 7864320  | 693,348           | 207601 | 34734080 | 619,896 | 207601 | 61603840 | 624,616 | 207601 |
| 8519680  | 620,666           | 207601 | 35389440 | 619,974 | 207601 | 62259200 | 611,618 | 207601 |
| 9175040  | 615,21            | 207601 | 36044800 | 616,235 | 207601 | 62914560 | 612,904 | 207601 |
| 9830400  | 617,025           | 207601 | 36700160 | 618,435 | 207601 | 63569920 | 613,765 | 207601 |
| 10485760 | 631,104           | 207601 | 37355520 | 619,777 | 207601 | 64225280 | 614,611 | 207601 |
| 11141120 | 625,269           | 207601 | 38010880 | 622,355 | 207601 | 64880640 | 626,191 | 207601 |
| 11796480 | 642,488           | 207601 | 38666240 | 627,719 | 207601 | 65536000 | 617,454 | 207601 |
| 12451840 | 620,376           | 207601 | 39321600 | 622,91  | 207601 |          |         |        |
| 13107200 | 621,717           | 207601 | 39976960 | 619,842 | 207601 |          |         |        |
| 13762560 | 614,705           | 207601 | 40632320 | 613,966 | 207601 |          |         |        |
| 14417920 | 620,185           | 207601 | 41287680 | 620,551 | 207601 |          |         |        |
| 15073280 | 629,077           | 207601 | 41943040 | 618,346 | 207601 |          |         |        |
| 15728640 | 636,728           | 207601 | 42598400 | 618,348 | 207601 |          |         |        |
| 16384000 | 621,122           | 207601 | 43253760 | 619,835 | 207601 |          |         |        |
| 17039360 | 619,136           | 207601 | 43909120 | 618,507 | 207601 |          |         |        |
| 17694720 | 621,099           | 207601 | 44564480 | 624,307 | 207601 |          |         |        |
| 18350080 | 622,029           | 207601 | 45219840 | 612,281 | 207601 |          |         |        |
| 19005440 | 620,783           | 207601 | 45875200 | 615,453 | 207601 |          |         |        |
| 19660800 | 627,175           | 207601 | 46530560 | 619,584 | 207601 |          |         |        |
| 20316160 | 619,848           | 207601 | 47185920 | 620,695 | 207601 |          |         |        |
| 20971520 | 619,362           | 207601 | 47841280 | 617,415 | 207601 |          |         |        |
| 21626880 | 635,624           | 207601 | 48496640 | 613,059 | 207601 |          |         |        |
| 22282240 | 626,874           | 207601 | 49152000 | 665,952 | 207601 |          |         |        |
| 22937600 | 620,564           | 207601 | 49807360 | 617,371 | 207601 |          |         |        |
| 23592960 | 618,962           | 207601 | 50462720 | 615,477 | 207601 |          |         |        |
| 24248320 | 617,738           | 207601 | 51118080 | 616,365 | 207601 |          |         |        |
| 24903680 | 622,434           | 207601 | 51773440 | 617,455 | 207601 |          |         |        |
| 25559040 | 624,642           | 207601 | 52428800 | 617,158 | 207601 |          |         |        |
| 26214400 | 622,59            | 207601 | 53084160 | 646,461 | 207601 |          |         |        |
| 26869760 | 653,44            | 207601 | 53739520 | 618,287 | 207601 |          |         |        |

Figure 36: Execution time of the randjump function in Experiment 2 using the fourth implementation on the cast\_info file with B varying from 65536 to 65536000 and j = 10000

| B          | Average time (ms) | Sum    |            |          |        |            |          |        |
|------------|-------------------|--------|------------|----------|--------|------------|----------|--------|
| 1310720000 | 629,55            | 207601 |            |          |        |            |          |        |
| 1317273600 | 479,256           | 207601 | 1599078400 | 104,9646 | 207601 | 1880883200 | 78,4138  | 207601 |
| 1323827200 | 497,638           | 207601 | 1605632000 | 136,5192 | 207601 | 1887436800 | 79,9592  | 207601 |
| 1330380800 | 493,008           | 207601 | 1612185600 | 130,1274 | 207601 | 1893990400 | 80,5882  | 207601 |
| 1336934400 | 478,86            | 207601 | 1618739200 | 103,0018 | 207601 | 1900544000 | 77,9722  | 207601 |
| 1343488000 | 474,302           | 207601 | 1625292800 | 85,7642  | 207601 | 1907097600 | 78,1846  | 207601 |
| 1350041600 | 525,352           | 207601 | 1631846400 | 88,3928  | 207601 | 1913651200 | 77,4034  | 207601 |
| 1356595200 | 474,83            | 207601 | 1638400000 | 86,024   | 207601 | 1920204800 | 81,1044  | 207601 |
| 1363148800 | 458,32            | 207601 | 1644953600 | 84,0608  | 207601 | 1926758400 | 80,4268  | 207601 |
| 1369702400 | 464,63            | 207601 | 1651507200 | 87,9208  | 207601 | 1933312000 | 78,8774  | 207601 |
| 1376256000 | 458,57            | 207601 | 1658060800 | 87,0684  | 207601 | 1939865600 | 82,3238  | 207601 |
| 1382809600 | 473,222           | 207601 | 1664614400 | 89,0218  | 207601 | 1946419200 | 79,8718  | 207601 |
| 1389363200 | 440,366           | 207601 | 1671168000 | 94,5058  | 207601 | 1952972800 | 79,5724  | 207601 |
| 1395916800 | 442,318           | 207601 | 1677721600 | 91,663   | 207601 | 1959526400 | 82,0276  | 207601 |
| 1402470400 | 489,738           | 207601 | 1684275200 | 90,0528  | 207601 | 1966080000 | 79,5492  | 207601 |
| 1409024000 | 504,948           | 207601 | 1690828800 | 83,054   | 207601 | 1972633600 | 80,665   | 207601 |
| 1415577600 | 409,188           | 207601 | 1697382400 | 80,3668  | 207601 | 1979187200 | 79,435   | 207601 |
| 1422131200 | 401,078           | 207601 | 1703936000 | 81,8678  | 207601 | 1985740800 | 81,1448  | 207601 |
| 1428684800 | 400,606           | 207601 | 1710489600 | 79,8408  | 207601 | 1992294400 | 79,6778  | 207601 |
| 1435238400 | 374,382           | 207601 | 1717043200 | 83,5156  | 207601 | 1998848000 | 78,765   | 207601 |
| 1441792000 | 352,866           | 207601 | 1723596800 | 88,2014  | 207601 | 2005401600 | 96,5054  | 207601 |
| 1448345600 | 312,412           | 207601 | 1730150400 | 92,6426  | 207601 | 2011955200 | 114,4146 | 207601 |
| 1454899200 | 114,6226          | 207601 | 1736704000 | 82,0496  | 207601 | 2018508800 | 82,7848  | 207601 |
| 1461452800 | 145,7226          | 207601 | 1743257600 | 81,0444  | 207601 | 2025062400 | 80,943   | 207601 |
| 1468006400 | 127,5082          | 207601 | 1749811200 | 79,3804  | 207601 | 2031616000 | 79,969   | 207601 |
| 1474560000 | 128,4316          | 207601 | 1756364800 | 79,276   | 207601 | 2038169600 | 79,3842  | 207601 |
| 1481113600 | 111,2634          | 207601 | 1762918400 | 80,9716  | 207601 | 2044723200 | 80,5236  | 207601 |
| 1487667200 | 88,2618           | 207601 | 1769472000 | 79,2068  | 207601 | 2051276800 | 81,6886  | 207601 |
| 1494220800 | 90,8296           | 207601 | 1776025600 | 86,544   | 207601 | 2057830400 | 80,1464  | 207601 |
| 1500774400 | 86,8736           | 207601 | 1782579200 | 81,6104  | 207601 | 2064384000 | 78,5156  | 207601 |
| 1507328000 | 85,815            | 207601 | 1789132800 | 82,5856  | 207601 | 2070937600 | 79,6138  | 207601 |
| 1513881600 | 89,8234           | 207601 | 1795686400 | 81,5602  | 207601 | 2077491200 | 79,3006  | 207601 |
| 1520435200 | 87,4112           | 207601 | 1802240000 | 79,7214  | 207601 | 2084044800 | 80,2886  | 207601 |
| 1526988800 | 89,728            | 207601 | 1808793600 | 79,217   | 207601 | 2090598400 | 78,8944  | 207601 |
| 1533542400 | 87,0782           | 207601 | 1815347200 | 77,728   | 207601 | 2097152000 | 80,4384  | 207601 |
| 1540096000 | 121,1008          | 207601 | 1821900800 | 82,3038  | 207601 | 2103705600 | 81,7884  | 207601 |
| 1546649600 | 138,617           | 207601 | 1828454400 | 79,2554  | 207601 | 2110259200 | 81,8084  | 207601 |
| 1553203200 | 97,6336           | 207601 | 1835008000 | 79,2592  | 207601 | 2116812800 | 83,2022  | 207601 |
| 1559756800 | 86,7128           | 207601 | 1841561600 | 80,4288  | 207601 | 2123366400 | 80,851   | 207601 |
| 1566310400 | 89,2328           | 207601 | 1848115200 | 79,0746  | 207601 | 2129920000 | 83,1116  | 207601 |
| 1572864000 | 86,0966           | 207601 | 1854668800 | 80,5002  | 207601 |            |          |        |
| 1579417600 | 89,5436           | 207601 | 1861222400 | 79,036   | 207601 |            |          |        |
| 1585971200 | 89,7538           | 207601 | 1867776000 | 78,8498  | 207601 |            |          |        |
| 1592524800 | 87,9948           | 207601 | 1874329600 | 79,534   | 207601 |            |          |        |

Figure 37: Execution time of the randjump function in Experiment 2 using the fourth implementation on the cast\_info file with B varying from 1310720000 to 2129920000 and j = 10000

## C Appendix 3 : Experiment 3

A = link\_type,kind\_type

B = link\_type,kind\_type,info\_type,company\_type,comp\_cast\_type,role\_type,movie\_link

C = complete\_cast,keyword

D = keyword,company\_name,movie\_info\_idx,movie\_keyword,movie\_companies,aka\_title,aka\_name

E = complete\_cast,keyword,movie\_link,company\_name,movie\_info\_idx,link\_type,kind\_type,info\_type,company\_type,comp\_cast\_type,role\_type,movie\_keyword,movie\_companies,aka\_title,aka\_name

F = movie\_info,person\_info

G = movie\_info,person\_info,name,cast\_info,title,movie\_keyword,char\_name

H = title,link\_type,kind\_type,info\_type,company\_type,comp\_cast\_type,role\_type,complete\_cast,keyword,movie\_link,company\_name,movie\_info\_idx,movie\_keyword,movie\_companies,aka\_title,aka\_name,movie\_info,person\_info,name,cast\_info,char\_name

The parameter k corresponds to the number of files.

- k=2 for A,C and F.
- k=7 for B,D and G.
- k=15 for E.
- k=21 for H.

| B \ files | A      | B       | C        | D       | E        | F      | G        | H        |
|-----------|--------|---------|----------|---------|----------|--------|----------|----------|
| 20        | 2.7984 | 348.364 | 3146.775 | 256671  | 200145.1 | 522209 | 1.85E+06 | 1.92E+06 |
| 40        | 2.4657 | 294.813 | 2755.22  | 216324  | 143275   | 419120 | 1.26E+06 | 1.37E+06 |
| 60        | 2.3554 | 227.318 | 2162.74  | 174769  | 111621   | 320816 | 979280   | 1.06E+06 |
| 80        | 2.1458 | 226.186 | 1908.02  | 147152  | 96014.5  | 275810 | 843360   | 908807   |
| 100       | 2.1457 | 194.339 | 1718.57  | 137235  | 88202.4  | 247398 | 751176   | 824213   |
| 120       | 2.1112 | 167.961 | 1645.79  | 121560  | 80430.5  | 225686 | 698587   | 757953   |
| 140       | 2.11   | 159.64  | 1559.54  | 118677  | 75121.7  | 209789 | 654416   | 707384   |
| 160       | 2.0894 | 151.396 | 1496.64  | 111405  | 73207.8  | 203786 | 626532   | 683284   |
| 180       | 2.1415 | 141.537 | 1488.5   | 105502  | 71589.9  | 191232 | 609867   | 646340   |
| 200       | 2.1201 | 136.335 | 1435.05  | 104887  | 67896.7  | 186192 | 586400   | 629764   |
| 400       | 2.1024 | 134.218 | 1235.71  | 86268.1 | 56791.7  | 152744 | 502703   | 532783   |
| 600       | 2.0962 | 133.534 | 1130.47  | 82904.8 | 53229.9  | 142879 | 472759   | 527115   |
| 800       | 2.0647 | 126.851 | 1105.62  | 78822.7 | 53019.2  | 141336 | 454349   | 503109   |
| 1000      | 1.9914 | 129.375 | 1084.42  | 53626.1 | 50309.6  | 143226 | 450683   | 483385   |
| 1200      | 1.881  | 114.843 | 1091.83  | 49635   | 50020.6  | 130313 | 465640   | 477155   |
| 1400      | 1.8721 | 115.6   | 1103.05  | 42868.8 | 48749.7  | 127930 | 445603   | 434547   |
| 1600      | 1.8724 | 116.34  | 1049.51  | 42499.2 | 48771    | 129024 | 488503   | 445476   |
| 1800      | 1.8765 | 115.037 | 1085.53  | 48802.7 | 49117.3  | 125876 | 446415   | 429899   |
| 2000      | 1.8721 | 113.727 | 1046.87  | 47096.9 | 49823.3  | 127123 | 427604   | 436578   |

Figure 38: Execution time (in ms) of the pair InputStream2,OutputStream3 in function of B and the files

| B\files | A      | B       | C       | D       | E       | F      | G        | H        |
|---------|--------|---------|---------|---------|---------|--------|----------|----------|
| 20      | 3.6504 | 395.331 | 33457.2 | 195421  | 201564  | 614423 | 1.91E+06 | 2.25E+06 |
| 40      | 3.3145 | 331.18  | 2966.46 | 146884  | 150095  | 478749 | 1.53E+06 | 1.60E+06 |
| 60      | 3.145  | 234.966 | 2821.5  | 118753  | 117191  | 361826 | 1.27E+06 | 1.27E+06 |
| 80      | 3.0676 | 210.051 | 2398.48 | 101332  | 103799  | 307751 | 1.11E+06 | 1.16E+06 |
| 100     | 2.7142 | 190.101 | 2242.12 | 91799.2 | 92148.4 | 271947 | 1.01E+06 | 1.07E+06 |
| 120     | 2.4546 | 185.61  | 2120.7  | 86847.2 | 85845.9 | 250892 | 929928   | 962718   |
| 140     | 2.3598 | 168.606 | 1749.87 | 80346.1 | 81042.4 | 231387 | 912935   | 927552   |
| 160     | 2.2156 | 170.554 | 1572.15 | 76699.3 | 79075.4 | 222253 | 835499   | 879393   |
| 180     | 2.1564 | 166.69  | 1510.28 | 73786.1 | 74972.3 | 211533 | 840833   | 835650   |
| 200     | 2.1486 | 159.365 | 1513.2  | 71114.3 | 72475.6 | 204641 | 804007   | 809620   |
| 400     | 2.1121 | 146.324 | 1328.82 | 61130.5 | 64467.2 | 177213 | 748194   | 794142   |
| 600     | 2.0456 | 134.616 | 1245.69 | 57477.2 | 58402.4 | 157371 | 689429   | 680293   |
| 800     | 2.0693 | 128.207 | 1230.22 | 57990.9 | 56542.5 | 149161 | 734013   | 668924   |
| 1000    | 2.0679 | 120.422 | 1215.6  | 55049.6 | 56222.5 | 143359 | 674732   | 651864   |
| 1200    | 2.0456 | 116.497 | 1219.62 | 53986.2 | 55254.9 | 141394 | 709783   | 677435   |
| 1400    | 2.0497 | 120.817 | 1220.39 | 53547.1 | 54769.8 | 140502 | 641006   | 707816   |
| 1600    | 2.0465 | 131.865 | 1173.98 | 54075.5 | 55811.4 | 137836 | 663139   | 676364   |
| 1800    | 2.0467 | 136.47  | 1179.86 | 52821.6 | 53948.5 | 138079 | 632629   | 638047   |
| 2000    | 2.0466 | 131.698 | 1198.47 | 54977   | 53551.9 | 136303 | 644177   | 623149   |

Figure 39: Execution time (in ms) of the pair InputStream4,OutputStream3 in function of B and the files

| B\files    | A          | B        | C          | D          | E          | F          | G          | H          |
|------------|------------|----------|------------|------------|------------|------------|------------|------------|
| 65536      | 2.6611     | 115.788  | 1079.83    | 42713      | 46597.7    | 111515     | 394951     | 407167     |
| 131072     | 2.5541     | 112.414  | 1035.78    | 42412      | 45931.5    | 110656     | 389953     | 406182     |
| 327680     | 2.361      | 112.388  | 1051.25    | 42723      | 45327.1    | 110312     | 386902     | 404163     |
| 655360     | 2.9615     | 112.598  | 1059.83    | 42609      | 45136.5    | 110315     | 396235     | 404235     |
| 1310720    | 2.4299     | 113.908  | 1029.41    | 42601      | 45428.5    | 110303     | 387999     | 403523     |
| 6553600    | 2.6691     | 113.481  | 1039.18    | 42819      | 45230.1    | 110288     | 386451     | 402530     |
| Mean value | 2.60611667 | 113.4295 | 1049.21333 | 42646.1667 | 45608.5667 | 110564.833 | 390415.167 | 404633.333 |

Figure 40: Execution time (in ms) of the pair InputStream2,OutputStream4 in function of B and the files

| B\files    | A      | B          | C          | D        | E        | F      | G          | H          |
|------------|--------|------------|------------|----------|----------|--------|------------|------------|
| 65536      | 2.9893 | 122.544    | 1172.24    | 48742.9  | 50545.3  | 120042 | 579235     | 590157     |
| 131072     | 2.8206 | 111.798    | 1162.99    | 48672.2  | 49069.5  | 119695 | 577355     | 584375     |
| 327680     | 2.8292 | 117.244    | 1161.67    | 48610.1  | 49051.5  | 118595 | 571764     | 579654     |
| 655360     | 2.9293 | 114.978    | 1155.21    | 48574.8  | 47890.4  | 118493 | 568900     | 575155     |
| 1310720    | 2.7293 | 119.43     | 1099.46    | 48582.9  | 47999.5  | 119007 | 571024     | 574111     |
| 6553600    | 3.1293 | 121.665    | 1112.54    | 48589.6  | 48163.1  | 118234 | 567644     | 574743     |
| Mean value | 2.9045 | 117.943167 | 1144.01833 | 48628.75 | 48786.55 | 119011 | 572653.667 | 579699.167 |

Figure 41: Execution time (in ms) of the pair InputStream4,OutputStream4 in function of B and the files

## D Appendix 4 : Multi-way Merge sort

| <u>M</u> | <u>Average time (t)</u> | <u>t/log(ceil(N/M))</u>   |
|----------|-------------------------|---------------------------|
| 10000    | 182307                  | 56083,00474               |
| 20000    | 158778                  | 53825,28224               |
| 30000    | 149701                  | 53969,90827               |
| 40000    | 145973                  | 55097,98036               |
| 50000    | 145741                  | 57093,59292               |
| 60000    | 143695                  | 58111,26285               |
| 70000    | 134600                  | 55930,91738               |
| 80000    | 133389                  | 56802,25004               |
| 90000    | 138139                  | 60147,64389               |
| 100000   | 134898                  | 59878,7396                |
| 110000   | 129970                  | 58822,86345               |
| 120000   | 127892                  | 58849,99453               |
| 130000   | 122551                  | 57354,71539               |
| 140000   | 124929                  | 59286,45071               |
| 150000   | 122407                  | 58975,77953               |
| 160000   | 126142                  | 61556,16367               |
| 170000   | 115869                  | 57327,13905               |
| 180000   | 108678                  | 54457,84896               |
| 190000   | 113753                  | 57651,10446               |
| 200000   | 112284                  | 57456,53339               |
| 210000   | 112833                  | 58480,30124               |
| 220000   | 108280                  | 56736,10164               |
| 230000   | 108324                  | 57250,83717               |
| 240000   | 109450                  | 58371,42612               |
| 250000   | 110546                  | 59518,69157               |
| 260000   | 109880                  | 59754,76756               |
| 270000   | 106751                  | 58669,09719               |
| 280000   | 102242                  | 56606,76205               |
| 290000   | 102864                  | 57389,24176               |
| 300000   | 97286,9                 | 54712,38736               |
|          | <b><u>AVG</u></b>       | <b><u>57538,95964</u></b> |
|          | <b><u>STD</u></b>       | <b><u>1880,165646</u></b> |
|          | <b><u>CV</u></b>        | <b><u>0,032676393</u></b> |

Figure 42: Execution time of the multi-way merge sort algorithm on company\_name in function of M

| <b>File name</b> | <b>File length (N)</b> | <b>Average time (t)</b> | <b>N.log(N)</b> | <b>N.log(ceil(N/M))/t</b> |
|------------------|------------------------|-------------------------|-----------------|---------------------------|
| aka_name         | 73004391               | 694323                  | 574059003,7     | 269,5364772               |
| aka_title        | 38858446               | 313140                  | 294915609,7     | 284,1770016               |
| company_name     | 17802021               | 114098,2                | 129073006,9     | 304,9081072               |
| complete_cast    | 2414495                | 21334,9                 | 15411302,21     | 126,0662414               |
| keyword          | 3791536                | 17931,8                 | 24943814,61     | 270,3822435               |
| movie_companies  | 93112112               | 1,03E+06                | 742010995,7     | 241,6399684               |
| movie_info_idx   | 35335875               | 365914                  | 266722811,7     | 217,0840753               |
| movie_keyword    | 93799307               | 1,01E+06                | 747786800,5     | 247,9138787               |
| movie_link       | 656584                 | 3454,79                 | 3819539,731     | 114,4217036               |
| title            | 203877939              | 2,05E+06                | 1694097278      | 299,6678875               |
|                  |                        |                         | <b>AVG</b>      | <b>237,5797584</b>        |
|                  |                        |                         | <b>STD</b>      | <b>63,87730709</b>        |
|                  |                        |                         | <b>CV</b>       | <b>0,26886679</b>         |

Figure 43: Execution time of the multi-way merge sort algorithm on several files in function of N

| <b>d</b> | <b>Average time (t)</b> |
|----------|-------------------------|
| 2        | 167956                  |
| 3        | 146952                  |
| 4        | 141550                  |
| 5        | 139705                  |
| 6        | 128514                  |
| 7        | 123042                  |
| 8        | 119560                  |
| 9        | 116540                  |
| 10       | 116257                  |
| 11       | 120090                  |
| 12       | 116000                  |
| 13       | 109457                  |
| 14       | 111429                  |
| 15       | 109108                  |
| 16       | 108690                  |
| 17       | 109983                  |
| 18       | 109119                  |
| 19       | 107175                  |
| 20       | 109627                  |

Figure 44: Execution time of the multi-way merge sort algorithm on company\_name in function of d